

# Zeroth-Order Optimization

MSc in Advanced Computer Science

Trinity Term 2026

Candidate number: TODO

Word count: TODO

# Contents

|          |                                                                 |           |
|----------|-----------------------------------------------------------------|-----------|
| <b>1</b> | <b>Related Work</b>                                             | <b>1</b>  |
| 1.1      | Zeroth-Order Optimization and Evolutionary Strategies . . . . . | 1         |
| 1.2      | Scalable Zeroth-Order Methods . . . . .                         | 2         |
| 1.3      | Nonconvex Optimization and Saddle-Point Escape . . . . .        | 3         |
| 1.4      | Accelerated Optimization . . . . .                              | 3         |
| 1.5      | Continuous-Time Interpretations of Acceleration . . . . .       | 3         |
| <b>2</b> | <b>Preliminaries</b>                                            | <b>4</b>  |
| 2.1      | Optimization . . . . .                                          | 4         |
| 2.2      | Zero-Order Oracles . . . . .                                    | 7         |
| 2.3      | The Low-Rank Trick . . . . .                                    | 9         |
| 2.4      | Probability Tools . . . . .                                     | 10        |
| 2.5      | Notation . . . . .                                              | 14        |
| <b>3</b> | <b>Zeroth-Order Mirror Descent</b>                              | <b>15</b> |
| 3.1      | Classic Mirror Descent . . . . .                                | 16        |
| 3.2      | Zero-Order Mirror Descent . . . . .                             | 18        |
| 3.3      | Regret Bound . . . . .                                          | 19        |

|          |                                                  |           |
|----------|--------------------------------------------------|-----------|
| 3.4      | Concentration Bounds . . . . .                   | 20        |
| 3.5      | Energy Bound . . . . .                           | 22        |
| 3.6      | Putting it all Together . . . . .                | 24        |
| <b>4</b> | <b>Accelerated Zeroth-Order Gradient Descent</b> | <b>27</b> |
| 4.1      | The Main Result . . . . .                        | 29        |
| 4.2      | Quadratic Error . . . . .                        | 30        |
| 4.3      | Gradient Progress . . . . .                      | 31        |
| 4.4      | Mirror Descent against Progress . . . . .        | 35        |
| 4.5      | Linear Mixing . . . . .                          | 37        |
| 4.6      | Dealing with the Martingale . . . . .            | 38        |
| 4.7      | Putting it all Together . . . . .                | 39        |
| <b>5</b> | <b>EGGROLL for Nonconvex Optimization</b>        | <b>45</b> |
| 5.1      | Main Result . . . . .                            | 46        |
| 5.2      | Concentration Bounds . . . . .                   | 49        |
| 5.3      | Function Decrease . . . . .                      | 57        |
| 5.4      | Gradients Don't Move too Quickly . . . . .       | 61        |
| 5.5      | Progress Lower Bound . . . . .                   | 64        |
| 5.6      | Putting it all Together . . . . .                | 67        |



# Chapter 1

## Related Work

### 1.1 Zeroth-Order Optimization and Evolutionary Strategies

The first evolutionary strategy originated in Rechenberg's experimental work on aerodynamic design. In a wind-tunnel experiment, he repeatedly applied small random mutations to the five joint angles of a segmented plane, retaining a mutation whenever it reduced measured drag. After only around 300 mutations, this process was able to find a minimum-drag configuration Rechenberg (1965). Later, he developed the theoretical basis behind the  $(1+1)$ -ES and the celebrated  $1/5$  success rule Rechenberg (1973). His was the first of many algorithms which are now known as zeroth-order optimization algorithms, who can query an oracle  $f(x)$  but do not have access to its gradient  $\nabla f(x)$ .

Evolutionary strategies have been developed further since then. Instead of just using one mutation per time step, one could use a larger population size. The question then becomes, how to use the recombine the function evaluations most effectively. This led to Hansen and Ostermeier (2001) developing Covariance Matrix Adaptation Evolution Strategy (CMA-ES) that adopts a Gaussian distribution over candidate solutions, gradually moving the mean

towards a weighted average of the best candidates and updating the covariance to align with directions that repeatedly produce improvement.

Next came Natural Evolutionary Strategies (NES) which reduces expected fitness using the score-function identity  $\nabla_{\theta} \mathbb{E}_{z \sim p_{\theta}}[f(z)] = \mathbb{E}_{z \sim p_{\theta}}[f(z) \nabla_{\theta} \log p_{\theta}(z)]$  Wierstra et al. (2014) then rescaling this gradient estimate using an approximate Fisher information preconditioner. The main idea is that the important geometry is induced by the KL divergence between two policies, not the raw gradient in  $\theta$ . While these methods demonstrated remarkable empirical performance, the problem is that this Fisher preconditioner and covariance matrix used in CMA-ES are  $d \times d$ , which is impractical for large ML models.

## 1.2 Scalable Zeroth-Order Methods

This gave rise to scalable zeroth-order methods. Salimans et al. (2017) directly optimize the smoothed surrogate  $f_{\sigma}(x) = \mathbb{E}_{p \sim \mathcal{N}(0, I)}[f(x + \sigma p)]$  of  $f$  whose derivative satisfies  $\nabla_x f_{\sigma}(x) = \frac{1}{\sigma} \mathbb{E}_{p \sim \mathcal{N}(0, I)}[f(x + \sigma p)p]$ . They then use the Monte-Carlo estimator with  $r$  parallel workers

$$\frac{1}{r} \sum_{i=1}^r \frac{f(x + \sigma p_i) - f(x - \sigma p_i)}{2\sigma} p_i$$

to repeatedly update the parameters thereby decreasing  $f_{\sigma}(x)$ . To scale this approach, instead of sharing full perturbations across workers, they use pseudo-random generators to sample perturbations, and then they only need to share the seeds between CPUs, allowing each GPU to reconstruct every perturbation locally. Marking the first time evolutionary strategies was scaled, they were able to use over 1000 parallel workers.

For large ML models with trillions of parameters, 1000 parallel workers is not nearly enough to optimize losses in any reasonable time. The central bottleneck is that, with  $r$  parallel workers, each needs to hold  $O(r \times d)$  parameters where  $d$  is the dimension the each Gaussian perturbations. This is simply impossible to scale to large population size as  $d$  is

too large. Sarkar et al. (2025) have the solution: instead of using full Gaussian perturbations  $P \sim \mathcal{N}(0, I_{d \times m})$ , we can use low-rank perturbations  $E = \frac{1}{\sqrt{k}}AB^\top$  where  $A \sim \mathcal{N}(0, I_{d \times k})$  and  $B \sim \mathcal{N}(0, I_{m \times k})$ . Then, we can exploit  $u(M + \sigma E)^\top = uM^\top + \frac{\sigma}{\sqrt{k}}(uB)A^\top$ . They were able to scale evolutionary strategies for the first time to millions of parallel workers, demonstrating that zeroth-order methods can be used in practice for large ML models.

### 1.3 Nonconvex Optimization and Saddle-Point Escape

#### 1.4 Accelerated Optimization

#### 1.5 Continuous-Time Interpretations of Acceleration

## Chapter 2

# Preliminaries

In this chapter we provide some background in optimization, list the standard assumptions and definitions, and introduce the relevant probabilistic tools that will become central to our analysis in the sequel.

### 2.1 Optimization

The perfect optimization algorithm would take a function  $f : \mathbb{R}^d \rightarrow \mathbb{R}$  and output a point  $x^* \in \mathbb{R}^d$  satisfying

$$x^* \in \operatorname{argmin}_{x \in \mathbb{R}^d} f(x)$$

However, since deciding whether a quartic polynomial  $p(x_1, \dots, x_d) \in \mathbb{R}[x_1, \dots, x_d]$  is globally non-negative,

$$p(x) \geq 0 \quad \text{for all } x \in \mathbb{R}^d$$

is NP-hard Ahmadi et al. (2013), it is proven in Anandkumar and Ge (2016) that even verifying whether a given point is a local minimum of a non-convex function is NP-hard



in the worst case. So, the best we can hope for must be more restrictive. For general non-convex functions, the right definition ends up being

**Definition 2.1.** A point  $x \in \mathbb{R}^d$  is an  $\varepsilon$ -first order stationary point if

$$\|\nabla f(x)\| \leq \varepsilon$$

On the other hand, for convex  $f$ , we usually measure accuracy in terms of the function gap.

**Definition 2.2.** Let  $x^* \in \operatorname{argmin}_{x \in \mathbb{R}^d} f(x)$ . A point  $x \in \mathbb{R}^d$  is  $\varepsilon$ -optimal if

$$f(x) - f(x^*) \leq \varepsilon$$

To reach a first-order stationary point, we need an idea of how much signal we get from the gradient versus noise. Recall that for any differentiable function  $f(x)$ ,

$$f(x+h) = f(x) + \langle \nabla f(x), h \rangle + O(\|h\|^2)$$

The noise  $O(\|h\|^2)$  term is quantified by the  $L$ -smooth assumption

**Definition 2.3.** A function  $f \in \mathcal{C}^1(\mathbb{R}^d)$  is  $L$ -smooth if

$$\|\nabla f(x) - \nabla f(y)\| \leq L \|x - y\|$$

Which leads to the descent lemma, the starting point to the analysis of any gradient-based algorithm.

**Lemma 2.4** (Descent Lemma). *If  $f \in \mathcal{C}^1(\mathbb{R}^d)$  is  $L$ -smooth, then for any  $h \in \mathbb{R}^d$ ,*

$$f(x+h) \leq f(x) + \langle \nabla f(x), h \rangle + \frac{L}{2} \|h\|^2$$

*Proof.* Writing  $g(t) = f(x+th)$ , we know that  $\frac{d}{dt}g(t) = \langle \nabla f(x+th), h \rangle$ , and therefore

$$g(1) - g(0) = \int_0^1 \frac{d}{dt}g(t)dt = \int_0^1 \langle \nabla f(x+th), h \rangle dt$$

Therefore,

$$f(x+h) - f(x) - \langle \nabla f(x), h \rangle = \int_0^1 \langle \nabla f(x+th) - \nabla f(x), h \rangle dt$$

By the Cauchy-Schwarz inequality and the  $L$ -smooth assumption,

$$\begin{aligned} |\langle \nabla f(x+th) - \nabla f(x), h \rangle| &\leq \|\nabla f(x+th) - \nabla f(x)\| \|h\| \\ &\leq tL \|h\|^2 \end{aligned}$$

Therefore,

$$|f(x+h) - f(x) - \langle \nabla f(x), h \rangle| \leq L \|h\|^2 \int_0^1 t dt = \frac{L}{2} \|h\|^2$$

□

The main optimization algorithm used in ML is gradient descent.

---

**Algorithm 1** Gradient Descent

---

**Require:** Objective function  $f : \mathbb{R}^d \rightarrow \mathbb{R}$ , smoothness constant  $L > 0$ , initial point  $x_0 \in \mathbb{R}^d$ ,  
number of iterations  $T$

- 1: **for**  $t = 0, \dots, T-1$  **do**
- 2:      $x_{t+1} \leftarrow x_t - \frac{1}{L} \nabla f(x_t)$
- 3: **end for**
- 4: **return**  $x_T$

---

By using the  $L$ -smoothness assumption, gradient descent converges to a first-order stationary point in  $O(L(f(x) - f^*)/\varepsilon^2)$  iterations where  $f^* = f(x^*) = \operatorname{argmin}_{x \in \mathbb{R}^d} f(x)$ . For convex functions, it demonstrates a linear convergence to an  $\varepsilon$ -optimal point in  $O(L\|x - x^*\|^2/\varepsilon)$  iterations. Gradient descent is the starting point for all algorithms discussed in the sequel.

## 2.2 Zero-Order Oracles

The focus of this work will be on zero-order optimization algorithms. Simply put, we drop the assumption that we can access gradients  $\nabla f(x)$ . We can only evaluate the function (“oracle”) at points of our choosing  $x \in \mathbb{R}^d$ ,  $f(x)$ . The simplest thing to do now is to calculate

$$\tilde{\nabla} f(x) = \sum_{i=1}^d \frac{f(x + \sigma e_i) - f(x - \sigma e_i)}{2\sigma} e_i$$

for any orthonormal basis  $\{e_1, \dots, e_d\}$  of  $\mathbb{R}^d$ . We have

$$\left\| \tilde{\nabla} f(x) - \nabla f(x) \right\| \leq \frac{L\sigma\sqrt{d}}{2}$$

This is a simple approach applied successfully, but is often of little practical use: frontier LLMs have  $d$  in the trillions, and trillions of forward passes to calculate just one gradient is far worse than the usual backpropagation.

For a quadratic  $f(x) = \frac{1}{2}x^\top Hx$  with symmetric  $H$ ,

$$\nabla f(x) = Hx = \left[ \frac{f(x+e_1)-f(x-e_1)}{2} \quad \dots \quad \frac{f(x+e_d)-f(x-e_d)}{2} \right]^\top$$

Thus typically the gradient gives  $d$  times more information than any single forward pass  $f(x + e_i)$ . This suggests that any zero-order optimization algorithm for convergence to a first-order stationary point for general nonconvex  $f$  should take  $d$  times more queries in the worst case. To the best of our knowledge this lower bound is currently open. Nevertheless, we are unaware of any method in the literature achieving a sublinear  $d$  dependence for this problem.

Another way to use the oracle  $f$  is to query random perturbations. For a random direction  $p \sim \mathcal{N}(0, I_d)$ , and search radius  $\sigma$ , we can look at  $f(x + \sigma p)$ . For  $p_1, \dots, p_r \stackrel{\text{i.i.d.}}{\sim} \mathcal{N}(0, I_d)$

where  $r$  is the population size, a Taylor expansion suggests the following estimator

$$\tilde{\nabla} f(x) = \frac{1}{r} \sum_{i=1}^r \frac{f(x + \sigma p_i) - f(x - \sigma p_i)}{2\sigma} p_i. \quad (2.1)$$

In the sequel, this is the only estimator for the gradient we use. Upon defining the gaussian density

$$\varphi_\sigma(z) = \frac{1}{(2\pi\sigma^2)^{d/2}} \exp\left(-\frac{\|z\|^2}{2}\right)$$

and the Gaussian smoothed objective

$$f_\sigma(x) = (f * \varphi_\sigma)(x) = \int_{\mathbb{R}^d} f(x + z) \varphi_\sigma(z) dz = \mathbb{E}_{p \sim \mathcal{N}(0, I_d)}[f(x + \sigma p)]$$

under standard integrability assumptions

$$\nabla f_\sigma(x) = \frac{1}{\sigma} \mathbb{E}[f(x + \sigma p)p]$$

which by symmetry of the Gaussian distribution yields,

$$\nabla f_\sigma(x) = \mathbb{E}[\tilde{\nabla} f(x)]$$

The advantage behind the central-difference version  $\tilde{\nabla} f(x)$  and  $\frac{1}{\sigma} f(x + \sigma p)p$  is that the first has good cancellation properties under a Taylor expansion

$$f(x + \sigma p) - f(x - \sigma p) = 2\sigma \langle \nabla f(x), p \rangle + O(\|\sigma p\|^3).$$

Since  $\mathbb{E}[\tilde{\nabla} f(x)] = \nabla f_\sigma(x)$ , the bias of the estimator is precisely the error introduced by Gaussian smoothing. The high-dimensional analysis of this bias was investigated in Sarkar et al. (2025), where they proved

$$|\mathbb{E}[\tilde{\nabla} f(x)] - \nabla f(x)| = O_d(L\sigma\sqrt{d})$$

Thus choosing  $\sigma = o(d^{-1/2})$  ensures this bias vanishes at infinity. The intuition is that  $f$  linearizes in high dimensions.

## 2.3 The Low-Rank Trick

In an ML model  $\Phi_\theta(x)$ , the parameters  $\theta \in \mathbb{R}^{d \times m}$  are naturally thought of not as vectors but as matrices. While helpful that the estimator 2.1 lets us evaluate at only  $r$  perturbations per iteration rather than the full  $d \times m$ , since  $p \in \mathbb{R}^{d \times m}$  has all Gaussian entries, every entry is non-zero almost surely. Thus evaluating  $\Phi_{\theta+p}(x)$  is still slow as we need to first add the two matrices  $\theta + p$  before evaluating, with memory  $O(rd)$  just to store the perturbations.

The low rank trick, introduced in Sarkar et al. (2025), is to sample  $a \sim \mathcal{N}(0, I_d)$ ,  $b \sim \mathcal{N}(0, I_m)$ , and exploit  $(\theta + ab^\top)x = \theta x + a(b^\top x)$ . This reduces the memory to  $O(r(d + m))$ , which makes zero-order algorithms practical.

However, for independent gaussian  $Z, W \sim \mathcal{N}(0, 1)$ , the product  $ZW$  is no longer subgaussian, it is sub-exponential, meaning it has fat tails. This could, *a priori*, ruin the convergence of the algorithm, as  $\mathbb{E}[\|\tilde{\nabla} f(x)\|^2]$  could grow too large. This is the central question of Chapter 5. To heed our analysis, for  $A, B \in \mathbb{R}^{d \times m}$ , we define the Frobenius inner product and the Frobenius norm,

$$\langle A, B \rangle_F = \text{Tr}(AB^\top) = \sum_{i,j} A_{i,j} B_{i,j}, \quad \|A\|_F = \sqrt{\sum_{i,j} A_{i,j}^2}$$

We define  $L$ -smoothness in the same way as before using this norm.

The Gaussian distribution is rotationally invariant, meaning for  $p \sim \mathcal{N}(0, I)$ , for any rotation matrix  $R$  satisfying  $RR^\top = R^\top R = I$ , we have  $Rp \sim \mathcal{N}(0, I)$ . Thus, we search all directions uniformly. However,  $ab^\top$  does not have this property, as it is rank 1, we could design a rotation matrix which rotates it into  $M$  which is full rank. This raises the question, is rotational invariance necessary? What if we miss the direction of the gradient?

**Definition 2.5.** The random variable  $X \in \mathbb{R}^d$  is isotropic if

$$\mathbb{E}[XX^\top] = I$$

this is equivalent to saying, for any  $u \in \mathbb{R}^d$ ,

$$\mathbb{E}[\langle X, u \rangle^2] = \mathbb{E}[u^\top X X^\top u] = u^\top \mathbb{E}[X X^\top] u = \|u\|^2$$

While  $ab^\top$  is not rotationally invariant, it is isotropic:

**Lemma 2.6.** *For  $a \sim \mathcal{N}(0, I_d)$ ,  $b \sim \mathcal{N}(0, I_m)$  independent, for any  $U \in \mathbb{R}^{d \times m}$ ,*

$$\mathbb{E}[\langle ab^\top, U \rangle_F^2] = \|U\|_F^2$$

*Proof.* We have

$$\mathbb{E}[\langle ab^\top, U \rangle_F^2] = \mathbb{E}[\text{Tr}(Uba^\top)^2]$$

By cyclic property of trace, we know that  $\text{Tr}(ABC) = \text{Tr}(CAB)$ , so the above equals

$$\mathbb{E}[(a^\top U b)^2] = \mathbb{E}[b^\top U^\top a a^\top U b]$$

By conditioning on  $b$ , using that  $a$  and  $b$  are independent, and that  $a \sim \mathcal{N}(0, I)$ ,

$$\mathbb{E}[b^\top \mathbb{E}[U^\top a a^\top U \mid b] b] = \mathbb{E}[b^\top U^\top \mathbb{E}[a a^\top] U b] = \mathbb{E}[b^\top U^\top U b]$$

This equals

$$\mathbb{E}[\text{Tr}(U^\top U b b^\top)] = \text{Tr}(U^\top U \mathbb{E}[b b^\top]) = \text{Tr}(U^\top U) = \|U\|_F^2.$$

□

It will turn out that isotropy is exactly the condition we need for convergence to a first-order stationary point.

## 2.4 Probability Tools

Here we give a short background on measure-theoretic probability. We define a  $\sigma$ -algebra

**Definition 2.7.** For a state space  $\Omega$ ,  $\mathcal{F} \subset 2^\Omega$  is a  $\sigma$ -algebra if:

1.  $\Omega \in \mathcal{F}$
2.  $\mathcal{F}$  is closed under complements: if  $A \in \mathcal{F}$ , then  $A^c \in \mathcal{F}$
3.  $\mathcal{F}$  is closed under countable unions: if  $A_1, A_2, \dots \in \mathcal{F}$ , then  $A = A_1 \cup A_2 \cup \dots \in \mathcal{F}$ .

By combining the last two properties  $\mathcal{F}$  is also closed under countable intersections.  $\mathcal{F}$  is the collection of *measurable sets*, or in probability, *events*.

**Definition 2.8.**  $\mathbb{P} : \mathcal{F} \rightarrow [0, 1]$  is a probability measure if  $\mathbb{P}(\emptyset) = 0$ ,  $\mathbb{P}(\Omega) = 1$ , and

$$\mathbb{P}\left(\bigcup_{i=0}^{\infty} E_i\right) = \sum_{i=0}^{\infty} \mathbb{P}(E_i)$$

The collection  $(\Omega, \mathcal{F}, \mathbb{P})$  is called a probability space.

In stochastic optimization, we often want to talk about the information available at iteration  $t$ . Filtrations give us a natural way to do this.

**Definition 2.9.**  $\{\mathcal{F}_t\}_{t=0}^{\infty}$  is a filtration on  $(\Omega, \mathcal{F}, \mathbb{P})$  if each  $\mathcal{F}_t \subset \mathcal{F}$  is  $\sigma$ -algebra, and

$$\mathcal{F}_0 \subset \mathcal{F}_1 \subset \mathcal{F}_2 \subset \dots$$

For Chapters 3 and 4, at each iteration  $t$  we sample  $p_{t,1}, \dots, p_{t,r} \sim \mathcal{N}(0, I_d)$  to evaluate  $\tilde{\nabla} f(x_t)$ . Thus we define the following filtration

**Definition 2.10.** The zero-order filtration  $\{\mathcal{F}_t\}_{t=0}^{\infty}$  is defined as

$$\mathcal{F}_t = \sigma(x_0, \dots, x_t, \{p_{\tau,i} : 0 \leq \tau \leq t-1, 1 \leq i \leq r\}).$$

which encapsulates all information available at time  $t$ .

Central to our analysis are concentration inequalities. Recall Hoeffding's inequality:

**Lemma 2.11** (Hoeffding's inequality). *Let  $X_1, \dots, X_n$  be independent random variables satisfying  $a_i \leq X_i \leq b_i$  almost surely for every  $1 \leq i \leq n$ . Then, for every  $t > 0$ ,*

$$\Pr\left(\left|\sum_{i=1}^n (X_i - \mathbb{E}[X_i])\right| \geq t\right) \leq 2 \exp\left(-\frac{2t^2}{\sum_{i=1}^n (b_i - a_i)^2}\right).$$

By re-arranging, we can get

**Lemma 2.12.** *Let  $X_1, \dots, X_n$  be independent random variables satisfying  $a_i \leq X_i \leq b_i$  almost surely for every  $1 \leq i \leq n$ . Then with probability at least  $1 - \delta$ , we have*

$$\left| \sum_{i=1}^n (X_i - \mathbb{E}[X_i]) \right| \leq \sqrt{\frac{1}{2} \sum_{i=1}^n (b_i - a_i)^2 \log(2/\delta)}$$

Using  $a_i \leq X_i \leq b_i$ ,  $\text{Var}(X_i) = \mathbb{E}[(X_i - \mathbb{E}(X_i))^2] \leq (b_i - a_i)^2$ . A bernoulli random variable which is  $a_i$  and  $b_i$  with probability  $1/2$  has variance  $(b_i - a_i)^2/4$ , so this is tight up to constant factors. Since  $X_i$  are independent,  $\text{Var}(X_1 + \dots + X_n) = \text{Var}(X_1) + \dots + \text{Var}(X_n) = \sum_{i=1}^n (b_i - a_i)^2$ . Hoeffding's inequality says that, when the  $a_i \leq X_i \leq b_i$ , the absolute value of the centered sum is not typically that much larger than its expectation.

However, this property is not limited to just  $a_i \leq X_i \leq b_i$ . It carries over to martingale difference sequences.

**Definition 2.13** (Martingale Difference Sequence).  $\{\xi_t\}_{t \geq 1}$  is a martingale difference sequence wrt filtration  $\mathcal{F}_t$  if  $\xi_t \in \mathcal{F}_t$  and

$$\mathbb{E}[\xi_t \mid \mathcal{F}_{t-1}] = 0$$

The main intuition is that in a martingale difference sequence, the variance of the full sum is the sum of the variance of each  $\xi_t$ , which is vastly superior than including all the cross covariance terms  $\text{Cov}(\xi_i, \xi_j)$ . This improves our bound substantially.

**Lemma 2.14.** *Define  $S_T = \sum_{t=1}^T \xi_t$ . Then*

$$\mathbb{E}[S_T^2] = \sum_{t=1}^T \mathbb{E}[\xi_t^2]$$

*Proof.* We have

$$\mathbb{E}[M_T^2] = \sum \mathbb{E}[\xi_t^2] + 2\mathbb{E}\left[\sum_{i < j} \xi_i \xi_j\right]$$



Now,

$$\mathbb{E} \left[ \sum_{i < j} \xi_i \xi_j \right] = \sum_{i < j} \mathbb{E}[\mathbb{E}[\xi_i \xi_j \mid \mathcal{F}_i]] = \sum_{i < j} \mathbb{E}[\xi_i \mathbb{E}[\xi_j \mid \mathcal{F}_i]] = 0.$$

□

For us, these martingale terms  $\xi_i$  will always be subgaussian.

**Definition 2.15.** The random variable  $X \in \mathbb{R}$  is  $\sigma$ -subgaussian if

$$\mathbb{E} \left[ \exp \left( \frac{X^2}{3\sigma^2} \right) \right] \leq 2$$

Under this definition  $\mathcal{N}(0, 1)$  is 1-subgaussian (the constant 3 is not standard). We have the following lemma which vastly generalizes Hoeffding's inequality and will help us substantially in the sequel,

**Lemma 2.16** (Adapted Sub-Gaussian Hoeffding's inequality). *Let  $\{\mathcal{F}_t\}_{t=0}^T$  be a filtration, and let  $X_t$  be  $\mathcal{F}_t$ -measurable. Suppose that, conditional on  $\mathcal{F}_{t-1}$ ,  $X_t$  is  $\sigma_t$ -subgaussian and mean zero. Then, for any  $\delta \in (0, 1)$ , with probability at least  $1 - \delta$ ,*

$$\left| \sum_{t=1}^T X_t \right| \leq c \sqrt{\log \left( \frac{2}{\delta} \right) \sum_{t=1}^T \sigma_t^2},$$

where  $c > 0$  is an absolute constant.

Convergence of stochastic gradient methods are standard in the literature, see for example Ghadimi and Lan (2013); Nesterov and Spokoiny (2017); Balasubramanian and Ghadimi (2018). Concentration inequalities will allow us to upgrade that convergence in expectation to convergence guarantees that hold with high probability.

## 2.5 Notation

$d$  and  $m$  will always be dimension constants,  $r$  will always be the population size,  $T$  will always be the iteration horizon,  $L$  will always be the smoothness constant,  $\sigma$  will always be the smoothing radius,  $\varepsilon$  will always be the desired accuracy, and  $\delta$  will always be the failure probability. We use  $O(\cdot)$ ,  $\Omega(\cdot)$ , and  $\Theta(\cdot)$  in the standard sense. The notation  $\tilde{O}(\cdot)$  hides polylogarithmic factors in the problem parameters, including  $d$ ,  $L$ ,  $T$ ,  $f(x_0) - f^*$ ,  $\varepsilon$  and failure probability  $\delta$ . Throughout the sequel,  $\chi$  always denotes a polylogarithmic factor.

## Chapter 3

# Zeroth-Order Mirror Descent

Gradient descent iteratively takes a step in the gradient direction, gradually reducing function value. It successively builds better and better upper bounds on  $f(x^*)$ . However, it ignores the dual problem of finding lower bounds on  $f(x^*)$ . For convex functions, we have for all  $u \in \mathbb{R}^d$ ,

$$f(u) \geq f(x) + \langle \nabla f(x), u - x \rangle$$

Thus, for some points  $\{x_t\}_{t=0}^{T-1}$ , we have

$$f(u) \geq \frac{1}{T} \sum_{t=0}^{T-1} f(x_t) + \frac{1}{T} \sum_{t=0}^{T-1} \langle \nabla f(x_t), u - x \rangle \quad (3.1)$$

Defining  $\bar{x} = \frac{1}{T} \sum_{t=0}^{T-1} x_t$ , convexity yields  $f(\bar{x}) \leq \frac{1}{T} \sum_{t=0}^{T-1} f(x_t)$ , and combining this with the above yields,

$$f(\bar{x}) - f(u) \leq \frac{1}{T} \sum_{t=0}^{T-1} \langle \nabla f(x_t), x_t - u \rangle =: R_T(u)$$

$R_T(u)$ , using terminology from online learning, is referred to as the regret. Consider a regularized version

$$\tilde{R}_T(u) = \frac{1}{T} \left( -\frac{w(u)}{\alpha} + \sum_{t=0}^{T-1} \langle \nabla f(x_t), x_t - u \rangle \right)$$

where  $\alpha > 0$  is a trade-off parameter and  $w(u)$  is a strongly convex regularizer. By taking the updates  $x_t = \operatorname{argmax}_u \tilde{R}_T(u)$ , we will tune  $\alpha$  to show that Mirror descent converges in  $T = O(\rho^2/\varepsilon^2)$  iterations where  $\rho^2$  is the average value of  $\|\nabla f(x_t)\|^2$  across iterations. This is because intuitively, the upper bound (3.1) becomes better when  $\frac{1}{T} \sum_{t=0}^{T-1} \|\nabla f(x_t)\|^2$  is small.

**Assumption 3.1.**  $f$  is  $\rho$ -Lipschitz continuous and convex over a closed convex set  $Q \subset \mathbb{R}^d$ .

**Definition 3.2.**  $w : Q \rightarrow \mathbb{R}^d$  is a DGF if it is 1-strongly convex w.r.t.  $\|\cdot\|$ , in symbols,  $w(y) \geq w(x) + \langle \nabla w(x), y - x \rangle + \frac{1}{2} \|y - x\|^2$ . The Bregman divergence is defined

$$V_x(y) = w(y) - \langle \nabla w(y), y - x \rangle - w(x)$$

We have  $V_x(x) = 0$  and  $V_x(y) \geq \frac{1}{2} \|x - y\|^2 \geq 0$ .

Classic examples are  $w(x) = \frac{1}{2} \|x\|_2^2$  for the  $\ell^2$  norm which gives  $V_x(y) = \frac{1}{2} \|y - x\|_2^2$ , and  $w(y) = \sum_i y_i \log(y_i)$  for the  $\ell^1$  norm with  $Q = \Delta$ , which gives KL divergence  $V_x(y) = \sum_i y_i \log(y_i/x_i) \geq \frac{1}{2} \|x - y\|_1^2$ .

### 3.1 Classic Mirror Descent

**Definition 3.3** (Mirror Descent Step). The mirror descent step with step length  $\alpha$  is

$$x^+ = \operatorname{Mirr}_x(\alpha \nabla f(x)) \quad \text{where} \quad \operatorname{Mirr}_x(\xi) = \operatorname{argmin}_{y \in Q} \{V_x(y) + \langle \xi, y - x \rangle\}.$$

Mirror descents core lemma is this:

**Lemma 3.4** (Allen-Zhu and Orecchia (2016)). *If  $x_{t+1} = \text{Mirr}_{x_t}(\alpha \nabla f(x_t))$ , then*

$$\forall u \in Q, \quad \alpha(f(x_t) - f(u)) \leq \alpha \langle \nabla f(x_t), x_t - u \rangle \leq \frac{\alpha^2}{2} \|\nabla f(x_t)\|^2 + V_{x_t}(u) - V_{x_{t+1}}(u).$$

*Proof.* Notice

$$\alpha \langle \nabla f(x_t), x_t - u \rangle = \langle \alpha \nabla f(x_t), x_t - x_{t+1} \rangle + \langle \alpha \nabla f(x_t), x_{t+1} - u \rangle$$

Since  $Q$  is convex and  $x_{t+1} = \text{Mirr}_{x_t}(\alpha \nabla f(x_t))$  is minimal, it satisfies the first order condition

$$\langle \nabla V_{x_t}(x_{t+1}) + \alpha \nabla f(x_t), u - x_{t+1} \rangle \geq 0 \quad \forall u \in Q$$

Therefore the above is at most

$$\langle \alpha \nabla f(x_t), x_t - x_{t+1} \rangle + \langle -\nabla V_{x_t}(x_{t+1}), x_{t+1} - u \rangle$$

Using the three-point identity

$$\langle -\nabla V_x(y), y - u \rangle = V_x(u) - V_y(u) - V_x(y)$$

we get the above equals

$$\begin{aligned} & \langle \alpha \nabla f(x_t), x_t - x_{t+1} \rangle + V_{x_t}(u) - V_{x_{t+1}}(u) - V_{x_t}(x_{t+1}) \\ & \leq \left( \langle \alpha \nabla f(x_t), x_t - x_{t+1} \rangle - \frac{1}{2} \|x_t - x_{t+1}\|^2 \right) + (V_{x_t}(u) - V_{x_{t+1}}(u)) \\ & \leq \frac{\alpha^2}{2} \|\nabla f(x_t)\|^2 + (V_{x_t}(u) - V_{x_{t+1}}(u)) \end{aligned}$$

After recalling that  $V_x(y) \geq \frac{1}{2} \|x - y\|^2$ , using Cauchy-Schwarz  $\langle a, b \rangle \leq \|a\| \cdot \|b\|$  and Young's inequality  $2ab \leq a^2 + b^2$  in that order.

The first inequality follows because by convexity  $f(u) \geq f(x_t) + \langle \nabla f(x_t), u - x_t \rangle$ .  $\square$

Using this, we can prove Mirror Descent converges.

**Lemma 3.5** (Allen-Zhu and Orecchia (2016)). *For  $\bar{x} = \frac{1}{T} \sum_{t=0}^{T-1} f(x_t)$ , if we choose*

$\alpha = \sqrt{2\Theta/\rho^2 T}$ , we have

$$f(\bar{x}) - f(x^*) \leq \rho \sqrt{\frac{2\Theta}{T}}$$

*Proof.* Telescoping, we get

$$\sum_{t=0}^{T-1} \alpha(f(x_t) - f(x^*)) \leq \sum_{t=0}^{T-1} \frac{\alpha^2}{2} \|\nabla f(x_t)\|^2 + V_{x_0}(x^*) - V_{x_T}(x^*)$$

Using convexity,  $f(\bar{x}) \leq 1/T \sum_{t=0}^{T-1} f(x_t)$ , and recalling that  $f$  is  $\rho$ -Lipschitz continuous, we have

$$f(\bar{x}) - f(x^*) \leq \frac{\alpha}{2} \rho^2 + \frac{1}{\alpha} \Theta,$$

where  $\Theta$  is any number satisfying  $\Theta \geq V_{x_0}(x^*)$ . Now by choosing  $\alpha = \sqrt{2\Theta/\rho^2 T}$  to balance the two terms, we are done.  $\square$

## 3.2 Zero-Order Mirror Descent

In this section, we assume that the norm  $\|\cdot\|$  is the  $\ell^2$  norm and  $Q = \mathbb{R}^d$ . We begin by introducing the zeroth-order gradient estimator.

**Definition 3.6** (Zero-Order Estimator). The Zero-Order gradient estimator is defined as follows

$$\tilde{\nabla} f(x) = \frac{f(x + \sigma p) - f(x - \sigma p)}{2\sigma} \cdot p$$

where  $p \sim \mathcal{N}(0, I_d)$ .

The main result of this chapter is the following.

**Theorem 3.7** (informal). *With iterates defined as  $x_{t+1} = \text{Mirr}_{x_t}(\alpha \tilde{\nabla} f(x_t))$ , where the*

perturbation sampled at time  $t$  is independent from the past, if  $f$  is quadratic and

$$\alpha = \tilde{O} \left( \sqrt{\frac{\Theta}{\rho^2 T d}} \right)$$

if we let  $\bar{x} = \frac{1}{T} \sum_{t=0}^{T-1} x_t$ , with probability at least  $1 - \delta$ ,

$$f(\bar{x}) - f^* \leq \tilde{O} \left( \sqrt{\frac{d\Theta\rho^2}{T}} \right).$$

This matches the convergence rate of Mirror descent presented in Allen-Zhu and Orecchia (2016), up to the zero-order  $d$  penalty and log factors.

### 3.3 Regret Bound

Now, notice that the only place in the proof of Lemma 3.4 where we used that the update is evaluated at the true gradient  $\nabla f(x_t)$  was to upper bound  $f(x_t) - f(u)$ . The gradient was not used in the proof, actually. So we get this lemma for free:

**Lemma 3.8.** *If  $x_{t+1} = \text{Mirr}_{x_t}(\alpha \tilde{\nabla} f(x_t))$ , then*

$$\forall u \in Q \quad \alpha \langle \tilde{\nabla} f(x_t), x_t - u \rangle \leq \frac{\alpha^2}{2} \left\| \tilde{\nabla} f(x_t) \right\|^2 + V_{x_t}(u) - V_{x_{t+1}}(u).$$

Notice now that for any  $u \in Q$ ,

$$\alpha(f(x_t) - f(u)) \leq \alpha \langle \nabla f(x_t), x_t - u \rangle = \alpha \langle \tilde{\nabla} f(x_t), x_t - u \rangle + \alpha \langle \nabla f(x_t) - \tilde{\nabla} f(x_t), x_t - u \rangle$$

For simplicity of exposition in this chapter only we will assume there is no quadratic error and  $\tilde{\nabla} f(x_t) = p_t p_t^\top \nabla f(x_t)$ . Thus the above equals,

$$\begin{aligned} & \alpha \langle \tilde{\nabla} f(x_t), x_t - u \rangle + \alpha \langle (I - p_t p_t^\top) \nabla f(x_t), x_t - u \rangle \\ & \leq \frac{\alpha^2}{2} \left\| \tilde{\nabla} f(x_t) \right\|^2 + V_{x_t}(u) - V_{x_{t+1}}(u) + \alpha \langle (I - p_t p_t^\top) \nabla f(x_t), x_t - u \rangle \end{aligned} \quad (3.2)$$

### 3.4 Concentration Bounds

**Lemma 3.9.** *With probability at least  $1 - \delta$ , we have for every  $1 \leq t \leq T$ ,*

$$\left\| p_t p_t^\top \nabla f(x_t) \right\| \leq c\sqrt{d} \log(cT/\delta) \|\nabla f(x_t)\|$$

*Proof.* By Lemma 5.7, who's proof is deferred to Chapter 5, the event

$$A_t = \left\{ \|p_t\| \leq \sqrt{3d \log(C/\delta)} \right\}$$

has probability at least  $1 - \delta$ . Now, as  $p_t^\top \nabla f(x_t) \sim \mathcal{N}(0, \|\nabla f(x_t)\|^2)$  conditional on  $\mathcal{F}_t$ , we know that

$$\mathbb{E} \left[ \exp \left( \frac{\|p_t p_t^\top \nabla f(x_t) \mathbb{1}_A\|^2}{t^2} \right) \right] \leq \mathbb{E} \left[ \exp \left( \frac{3d \log(C/\delta) |p_t^\top \nabla f(x_t)|^2}{t^2} \right) \right]$$

Using  $\mathbb{E} \left[ e^{\lambda Z^2} \right] = (1 - 2\lambda\sigma^2)^{-1/2}$  for  $Z \sim \mathcal{N}(0, \sigma^2)$ , we know that

$$\mathbb{E} \left[ \exp \left( \frac{3d \log(C/\delta) |p_t^\top \nabla f(x_t)|^2}{t^2} \right) \mid \mathcal{F}_t \right] = \left( 1 - \frac{6d \log(C/\delta) \|\nabla f(x_t)\|^2}{t^2} \right)^{-1/2}$$

The RHS is  $\leq 2$  for  $t = \sqrt{8d \log(C/\delta)} \|\nabla f(x_t)\|$ . By taking expectations on both sides this bound also holds unconditionally. Then using Markov's inequality,

$$\begin{aligned} & \mathbb{P} \left( \left\| p_t p_t^\top \nabla f(x_t) \mathbb{1}_A \right\| \geq \sqrt{8d \log(C/\delta)} \|\nabla f(x_t)\| \lambda \right) \\ &= \mathbb{P} \left( \exp \left( \frac{\|p_t p_t^\top \nabla f(x_t) \mathbb{1}_A\|^2}{8d \log(C/\delta) \|\nabla f(x_t)\|^2} \right) \geq e^{\lambda^2} \right) \leq 2e^{-\lambda^2} \end{aligned}$$

Now choosing  $\lambda = \sqrt{\log(C/\delta)}$ , possibly changing  $C$  to ensure  $C \geq 2$ , we have, with probability at least  $1 - \delta$  on  $A$ ,

$$\left\| p_t p_t^\top \nabla f(x_t) \right\| \leq \log(C/\delta) \sqrt{8d} \|\nabla f(x_t)\|$$

Since  $A$  has probability at least  $1 - \delta$ , this bound holds with probability at least  $1 - 2\delta$ .

Rescaling  $\delta \leftarrow \delta/2$  and changing the absolute constant yields the lemma.  $\square$



This lemma highlights the main technique for creating concentration bounds when the underlying distribution has fat tails. The first thing you do is truncate with high probability, and prove the resulting random variable has similar scale (whether that be sub-gaussian, sub-exponential, etc.) as what you want, up to log factors.

The main idea behind the rest of this proof is that the mean-0 martingale error

$$\sum_{t=0}^{T-1} \alpha \langle (I - p_t p_t^\top) \nabla f(x_t), x_t - u \rangle$$

can hopefully be bounded by the other terms of the equation. In this case, it ends up being the “energy” term  $V_{x_t}(u)$  and the Lipschitz constant  $\rho$ .

The main lemma which is different from the normal mirror descent bounds is the following.

**Lemma 3.10.** *Let  $\{\mathcal{F}_t\}$  be a filtration. Assume that  $g_t, d_t \in \mathcal{F}_{t-1}$ , and assume  $p_t \sim \mathcal{N}(0, I)$  with  $p_t \in \mathcal{F}_t$ ,  $p_t \perp \mathcal{F}_{t-1}$ . Define*

$$E_\tau = \sum_{t=0}^{\tau-1} g_t^\top (I - p_t p_t^\top) d_t$$

*Then for any  $\theta$ ,  $0 < \delta < 1/e$ , with probability at least  $1 - \delta$ ,*

$$|E_\tau| \leq c\theta \sum_{t=0}^{\tau-1} \log\left(\frac{c\tau}{\delta}\right) \|g_t\|^2 \|d_t\|^2 + \frac{1}{\theta} \log\left(\frac{c\tau}{\delta}\right).$$

*Proof.* For each  $t$  let

$$A_t = \left\{ L_\delta \leq |p_t^\top d_t| \leq \sqrt{3 \log(c/\delta) \|d_t\|^2} \right\}, \quad \mathbb{P}_{\mathcal{F}_{t-1}}(A_t) \geq 1 - \delta$$

Then by a lemma from before we know that  $\mathbb{E}_{\mathcal{F}_{t-1}}[(I - p_t p_t^\top) d_t; A_t] = 0$ . In particular,

$$\mathbb{E}_{\mathcal{F}_{t-1}}[g_t^\top (I - p_t p_t^\top) d_t; A_t] = 0.$$

Dropping the subscript  $t$  for simplicity,

$$\mathbb{E} \left[ \exp \left( \frac{(g^\top p p^\top d \mathbb{1}_A)^2}{t^2} \right) \right] \leq \mathbb{E} \left[ \exp \left( \frac{3 \log(C/\delta) \|d\|^2 (p^\top g)^2}{t^2} \right) \right]$$

Now,  $p^\top g$  is  $\mathcal{N}(0, \|g\|^2)$  random variable. Therefore,  $\mathbb{E}\left[\exp\left((p^\top g)^2/4\|g\|^2\right)\right] \leq 2$ . Now, taking  $c \geq 1$  and  $0 < \delta \leq 1/e$ , we have

$$(g^\top (I - pp^\top) d)^2 \leq 2(g^\top d)^2 + 2(g^\top pp^\top d)^2$$

So, the random variable  $Q_t = g_t^\top (I - p_t p_t^\top) d_t \mathbb{1}_{A_t}$  is  $24 \|g_t\|^2 \|d_t\|^2 \log(C/\delta)$  subgaussian and mean 0 given  $\mathcal{F}_t$ . Thus by the (subgaussian) Hoeffding inequality with probability at least  $1 - \delta$  for any  $\theta$ ,

$$\left| \sum_{t=0}^{\tau-1} Q_t \right| \leq c\theta \sum_{t=0}^{\tau-1} \log(c/\delta) \|g_t\|^2 \|d_t\|^2 + \frac{1}{\theta} \log\left(\frac{c}{\delta}\right)$$

Now  $\sum_{t=0}^{\tau-1} Q_t = E_\tau$  on  $\bigcap A_t$ . Notice, since  $\mathbb{P}(A_t) \geq 1 - \delta$ ,  $\mathbb{P}(\bigcap A_t) \geq 1 - \tau\delta$ . Therefore with probability at least  $1 - (\tau + 1)\delta$ ,  $E_\tau$  satisfies the same bound as above. Rescaling  $\delta$  to  $\delta/(\tau + 1)$  and we have completed the proof.  $\square$

This above lemma is the key for making the argument work. Because  $V_x(y) \geq \frac{1}{2} \|x - y\|^2$ , and recalling that  $f$  is  $\rho$ -Lipschitz continuous, we have

$$\left| \sum_{t=0}^{\tau-1} \alpha \langle (I - p_t p_t^\top) \nabla f(x_t), x_t - u \rangle \right| \leq \alpha \sqrt{c \sum_{t=0}^{\tau-1} \log^2(c\tau/\delta) \rho^2 V_{x_t}(u)}$$

after rescaling  $c$  and by choosing the optimal  $\theta$  according to AM-GM  $(x + y)/2 \geq \sqrt{xy}$ . So, it remains to bound  $\sum V_{x_t}(u)$ . We can do this because Lemma 3.5 actually tells us a bound on  $V_{x_\tau}(x^*)$ .

### 3.5 Energy Bound

So, by telescoping Equation (3.2), we arrive at the following lemma.

**Lemma 3.11.** *Define  $\chi = \log(cT^2/\delta)$ . Then with probability at least  $1 - \delta$ , for every*

$$1 \leq \tau \leq T,$$

$$V_{x_\tau}(x^*) \leq \frac{c\tau d\chi^2\alpha^2}{2}\rho^2 + V_{x_0}(x^*) + \alpha\rho\sqrt{c\sum_{t=0}^{\tau-1}\chi^2V_{x_t}(x^*)}$$

*Proof.* Telescoping Equation (3.2), we get

$$\begin{aligned} \sum_{t=0}^{\tau-1} \alpha(f(x_t) - f(x^*)) &\leq \sum_{t=0}^{\tau-1} \frac{\alpha^2}{2} \left\| \tilde{\nabla} f(x_t) \right\|^2 + V_{x_0}(x^*) - V_{x_\tau}(x^*) \\ &\quad + \sum_{t=0}^{\tau-1} \alpha \langle (I - p_t p_t^\top) \nabla f(x_t), x_t - x^* \rangle \end{aligned}$$

Therefore, by the previous discussion, by Lemma 3.9 and using that  $f(x_t) \geq f(x^*)$  for any  $x_t$ ,

$$V_{x_\tau}(x^*) \leq \frac{c\tau d\chi^2\alpha^2}{2}\rho^2 + V_{x_0}(x^*) + \alpha\rho\sqrt{c\sum_{t=0}^{\tau-1} \log^2(c\tau/\delta)V_{x_t}(x^*)}$$

We can have this hold jointly for all  $1 \leq \tau \leq T$  by rescaling  $\delta \leftarrow \delta/T$ . Then by using  $\tau \leq T$ , we arrive at the statement of the lemma.  $\square$

Naturally, this above lemma should tell us how much  $V_{x_t}(x^*)$  can grow. This next lemma tells us an upper bound on how much that is.

**Lemma 3.12.** *Let  $\Theta$  be any upper bound on  $V_{x_0}(x^*)$ . Then for some absolute constant  $c$ ,*

$$\sup_{0 \leq t \leq T} V_{x_t}(x^*) \leq 2\Theta + cd\alpha^2\rho^2\chi^2T$$

*Proof.* Define  $M = \sup_{0 \leq t \leq T} V_{x_t}(x^*)$ . Notice that the right hand side is increasing in  $\tau$ .

Therefore applying sup to both sides, we get

$$\begin{aligned} M &\leq \frac{cTd\chi^2\alpha^2}{2}\rho^2 + \Theta + \alpha\rho\sqrt{c\sum_{t=0}^{T-1}\chi^2V_{x_t}(x^*)} \\ &\leq \frac{cTd\chi^2\alpha^2}{2}\rho^2 + \Theta + \alpha\rho\sqrt{cT\chi^2M} \end{aligned}$$

Using Young's inequality, we know that

$$\alpha\rho\sqrt{cT\chi^2M} \leq \frac{1}{2}M + c\alpha^2\rho^2\chi^2T$$

Therefore,

$$\begin{aligned} M &\leq cTd\chi^2\alpha^2\rho^2 + 2\Theta + 2c\alpha^2\rho^2\chi^2T \\ &\leq 2\Theta + cd\alpha^2\rho^2\chi^2T \end{aligned}$$

for some absolute constant  $c$  as desired.  $\square$

A natural question is: Is this a good bound at all? Shouldn't  $x_t$  be getting closer to the minimizer  $x^*$ ? But we showed  $V_{x_t}(x^*)$  grows linearly?

To answer this, recall from before that  $\alpha = \sqrt{2\Theta/\rho^2T}$ . Ignoring the polylog factor  $\chi^{3/2}$ , and dimension  $d$ , we have actually shown that  $V_{x_t}(x^*)$  stays bounded. But still, shouldn't it be decreasing to 0? Actually, it should not, since  $f$  could have many minimizers. The best bound you could hope to satisfy for every  $f$  is that it is just bounded.

For a similar choice of  $\alpha$ , this tells us the accumulated error

$$\sum_{t=0}^{T-1} \alpha \langle (I - p_t p_t^\top) \nabla f(x_t), x_t - x^* \rangle$$

is around the same size as  $V_{x_0}(x^*)$ . This observation is what let's us prove convergence of Zero-Order mirror descent.

### 3.6 Putting it all Together

**Theorem 3.13** (Zero-Order Mirror Descent Convergence). *Let  $\bar{x} = 1/T \sum_{t=0}^{T-1} x_t$ , if we choose*

$$\alpha = \sqrt{\frac{2\Theta}{\rho^2 c \chi^2 T d}}$$

then with probability at least  $1 - \delta$ ,

$$f(\bar{x}) - f(x^*) \leq \frac{1}{T\alpha} \cdot 9\Theta = \tilde{O}\left(\rho\sqrt{\frac{\Theta d}{T}}\right)$$

*Proof.* By telescoping and using  $V_{x_T}(x^*) \geq 0$ , we get

$$\sum_{t=0}^{T-1} \alpha(f(x_t) - f^*) \leq \frac{cTd\chi^2\alpha^2}{2}\rho^2 + \Theta + \alpha\sqrt{c\sum_{t=0}^{T-1}\chi^2\rho^2V_{x_t}(x^*)}$$

By using the previous lemma, we know that

$$\alpha\sqrt{c\sum_{t=0}^{T-1}\chi^2\rho^2V_{x_t}(x^*)} \leq \alpha\sqrt{cT\chi^2\rho^2(2\Theta + c\alpha^2\rho^2\chi^2T)}$$

By choosing  $\alpha = \sqrt{2\Theta/\rho^2c\chi^2Td} \leq \sqrt{2\Theta/\rho^2c\chi^2T}$  as  $d \geq 1$ , we know that

$$c\alpha^2\rho^2\chi^2T \leq 2\Theta$$

and therefore,

$$\alpha\sqrt{cT\chi^2\rho^2(2\Theta + c\alpha^2\rho^2\chi^2T)} \leq 4\Theta$$

Now again by our choice of  $\alpha$ ,

$$\frac{cTd\chi^2\alpha^2}{2}\rho^2 \leq 4\Theta$$

Therefore,

$$\alpha(f(\bar{x}) - f^*) \leq \frac{9\Theta}{T}$$

We conclude,

$$f(\bar{x}) - f^* \leq \frac{9\Theta}{T\alpha} = \rho c \sqrt{\frac{9}{2}} \sqrt{\frac{d\Theta}{T}} \chi = \tilde{O}\left(\rho\sqrt{\frac{d\Theta}{T}}\right).$$

□

The extension of this analysis to the setting without quadratic error is left to future work. The main purpose of this chapter was not merely to show that mirror descent

converges, but to determine whether zeroth-order mirror descent can be analyzed with the zeroth-order gradient estimator. Having answered this question affirmatively, the techniques developed here will be applied extensively in the next chapter.

## Chapter 4

# Accelerated Zeroth-Order Gradient Descent via Linear Mixing

Accelerated gradient descent for  $L$ -smooth convex functions was developed in Nesterov (1983), improving the iteration complexity of finding an  $\varepsilon$ -optimal point from  $O(1/\varepsilon)$  with gradient descent to  $O(1/\sqrt{\varepsilon})$ . The algorithm can be defined inductively as

$$\begin{aligned}x_k &= y_{k-1} - s\nabla f(y_{k-1}) \\y_k &= x_k + \frac{k-1}{k+2}(x_k - x_{k-1})\end{aligned}$$

The point  $y_k$  “looks ahead” in the direction of velocity  $x_k - x_{k-1}$  with coefficient  $(k-1)/(k+2) \approx 1 - 3/k$ , then  $x_{k+1}$  subsequently does a gradient step at that point. This is slightly different from the momentum method

$$\begin{aligned}x_{k+1} &= y_k - s\nabla f(x_k) \\y_k &= x_k + \beta(x_k - x_{k-1})\end{aligned}$$

which exhibits acceleration for strongly convex quadratics, but may not even converge for general convex  $f$ . It is a mystery why the precise coefficient  $(k-1)/(k+2)$  was chosen by Nesterov, except as an algebraic trick to make an energy telescope. Recently, by taking a limit as the step size  $s \rightarrow 0$ , Su et al. (2016) find that Nesterov's AGD yields the second-order scheme

$$\ddot{X}_t + \frac{3}{t}\dot{X}_t + \nabla f(X_t) = 0$$

This is unusual not least because Nesterov's AGD is a first-order method, only relying on gradients. Perhaps surprisingly, their analysis found that  $3 = 2 + 1$  the smallest constant which preserves acceleration. Unfortunately, these Lyapunov proofs prove difficult to analyze with high probability, as martingale fluctuations yield an additive term, so the energy could grow linearly rather than decreasing.

To our rescue is the linear mixing framework discovered by Allen-Zhu and Orecchia (2016). By “mixing” mirror and gradient descent,

$$\begin{aligned} x_{k+1} &= \lambda z_k + (1 - \lambda)y_k \\ y_k &= x_k - \frac{1}{L}\nabla f(x_k) \\ z_k &= \text{Mirr}_{z_{k-1}}(\alpha \nabla f(x_k)) \end{aligned}$$

after carefully tuning  $\alpha$  and  $\lambda$ , Allen-Zhu and Orecchia (2016) show this scheme can also achieve acceleration. The point  $x_{k+1}$  is a linear combination of  $y_k$  and  $z_k$ ,  $y_k$  does a gradient step at  $x_k$ , and  $z_k$  does a mirror step with  $\nabla f(x_k)$ . The idea is that, while mirror and gradient descent tell us to go to different places, we can use part of both answer to achieve a faster scheme than regular gradient descent. Most importantly for us, they do not rely on carefully engineered Lyapunov functions, and the analysis is relatively direct. This is what will allow us to design a zero-order accelerated gradient descent scheme.



## 4.1 The Main Result

**Assumption 4.1.** We assume that the smoothness constant of  $f$  is  $L \geq 1$ . Otherwise, we can replace it with 1. Also assume that  $0 < \delta < 1/e$ .

We use the linear mixing framework of Allen-Zhu and Orecchia (2016) to design an accelerated zero-order gradient method. At each iteration, the gradient sequence  $(y_t)$  and the mirror sequence  $(z_t)$  are mixed to form a common query point. The same zeroth-order estimate at this point is then used for both updates. The full procedure is given in Algorithm 2.

**Definition 4.2** (Rank- $r$  Symmetric Zeroth-Order Estimator). For  $x \in \mathbb{R}^d$  and independent directions  $p_1, \dots, p_r \sim \mathcal{N}(0, I_d)$ , define

$$\tilde{\nabla} f(x) := \frac{1}{r} \sum_{i=1}^r \frac{f(x + \sigma p_i) - f(x - \sigma p_i)}{2\sigma} p_i.$$

with each iteration using  $2r$  function evaluations.

---

**Algorithm 2** Rank- $r$  Zeroth-Order Linear Mixing

---

**Require:** Zeroth-order oracle  $f$  and initial states  $x_0, y_0, z_0 \in \mathbb{R}^d$

**Require:** Horizon  $T$ , rank  $r$ , stepsizes  $\eta, \alpha > 0$ , smoothing radius  $\sigma > 0$ , and mixing weight  $\lambda \in (0, 1)$

**for**  $t = 0, \dots, T - 1$  **do**

    Set  $x_{t+1} \leftarrow \lambda z_t + (1 - \lambda)y_t$

    Sample  $p_{t+1,1}, \dots, p_{t+1,r} \stackrel{\text{i.i.d.}}{\sim} \mathcal{N}(0, I_d)$  to compute  $\tilde{\nabla} f(x_{t+1})$

    Update  $y_{t+1} \leftarrow x_{t+1} - \eta \tilde{\nabla} f(x_{t+1})$

    Update  $z_{t+1} \leftarrow \text{Mirr}_{z_t}(\alpha \tilde{\nabla} f(x_{t+1}))$

**end for**

**return**  $\bar{x}_T \leftarrow \frac{1}{T} \sum_{t=1}^T x_t$

---

The main result of this chapter is the following (informal) theorem.

**Theorem 4.3** (informal). *For a given initial point  $x_0$  satisfying  $f(x_0) - f^* \leq \Delta$  and letting  $\Theta \geq V_{x_0}(x^*)$  be any upper bound, by choosing*

$$\eta = \tilde{O}\left(\frac{r}{dL}\right), \quad \alpha = \tilde{O}\left(\frac{r}{d}\sqrt{\frac{\Theta}{L\Delta}}\right), \quad \sigma^2 \leq \tilde{O}\left(\frac{r^3\Delta^{3/2}}{L^{7/2}d^6\sqrt{\Theta}}\right),$$

$$\frac{1-\lambda}{\lambda} = \tilde{O}\left(\frac{d^2L\Theta}{\Delta}\right), \quad r = \tilde{O}(1)$$

*Algorithm 2 outputs a point  $\bar{x}$  satisfying  $f(\bar{x}) - f^* \leq \Delta/2$  in*

$$T = \tilde{O}\left(\frac{d}{r}\sqrt{\frac{L\Theta}{\Delta}}\right)$$

*iterations.*

*Furthermore, by restarting Algorithm 2  $\log_2(\Delta/\varepsilon)$  times, we receive a point  $\bar{x}$  satisfying  $f(\bar{x}) - f^* \leq \varepsilon$  in*

$$T = \tilde{O}\left(\frac{d}{r}\sqrt{\frac{L\Theta}{\varepsilon}}\right)$$

*iterations.*

## 4.2 Quadratic Error

The first thing to do is to control the quadratic error in this estimator.

**Lemma 4.4** (Quadratic Error).

$$\left\| \tilde{\nabla} f(x) - \frac{1}{r} \sum_{i=1}^r p_i p_i^\top \nabla f(x) \right\| \leq L\sigma \frac{1}{r} \sum_{i=1}^r \|p_i\|^3$$

*Proof.* We have

$$\begin{aligned} \frac{f(x + \sigma p) - f(x - \sigma p)}{2\sigma} p - pp^\top \nabla f(x) &= \frac{1}{2\sigma} \left[ \int_{-1}^1 \frac{d}{dt} f(x + t\sigma p) - \sigma pp^\top \nabla f(x) dt \right] \\ &= \frac{1}{2\sigma} \left[ \int_{-1}^1 \langle \nabla f(x + t\sigma p) - \nabla f(x), \sigma p \rangle p dt \right] \end{aligned}$$

Now, by  $L$ -smoothness,

$$\left| \int_{-1}^1 \langle \nabla f(x + \sigma t p) - \nabla f(x), \sigma p \rangle p dt \right| \leq \int_{-1}^1 L \sigma^2 |t| \|p\|^3 dt$$

Therefore,

$$\left| \frac{f(x + \sigma p) - f(x - \sigma p)}{2\sigma} p - p p^\top \nabla f(x) \right| \leq L \sigma \|p\|^3$$

Then by triangle inequality,

$$\left\| \tilde{\nabla} f(x) - \frac{1}{r} \sum_{i=1}^r p_i p_i^\top \nabla f(x) \right\| \leq L \sigma \frac{1}{r} \sum_{i=1}^r \|p_i\|^3.$$

□

Unfortunately, this bound is necessary. Ideally we would like to just say that  $f(x + \sigma p) - f(x - \sigma p)/2\sigma p = p p^\top \nabla g(x)$  for some function  $g(x)$  related to  $f(x)$ . Possibly a smoothed variant  $f_\sigma$ . But this cannot possibly hold for every  $p$  unless  $f$  is quadratic in a neighborhood of  $x$ . To see this, we can use MVT to expand the LHS as  $\langle \nabla f(x + \theta \sigma p), p \rangle p$ . By taking  $p = e_i$ , we see that  $e_i^\top \nabla g(x) = e_i^\top \nabla f(x + \theta_i \sigma e_i)$ . Sending  $\sigma \rightarrow 0$ ,  $\nabla f(x + \theta_i \sigma e_i) \rightarrow \nabla f(x)$  under  $L$ -smoothness assumption, meaning that  $\nabla g(x) = \nabla f(x)$ . This would say that  $f$  is quadratic, which is not true in general.

### 4.3 Gradient Progress

These lemmas tell us how much progress the gradient descent step is making.

**Lemma 4.5** (Vershynin (2018) Corollary 7.3.2). *Let  $P \in \mathbb{R}^{d \times r}$  be a matrix with i.i.d. Gaussian entries. Define  $\|A\|_{\text{op}} = \sup_{u \in \mathbb{S}^{r-1}} \|Au\|_2$ . Then*

$$\mathbb{P}\left(\|P\|_{\text{op}} \geq \sqrt{d} + \sqrt{r} + s\right) \leq 2e^{-cs^2}$$

**Proposition 4.6.** *Define  $\chi = \log(CrT/\delta)$  and let  $P_t = \begin{bmatrix} p_{t,1} & \cdots & p_{t,r} \end{bmatrix}$ . Define the*

bounded operator norm event

$$\mathfrak{D}_t = \left\{ \frac{1}{\sqrt{r}} \|P\|_{\text{op}} \leq C \frac{d}{r} \sqrt{\log(cTr/\delta)} \right\}$$

Then on  $\mathfrak{D}_t$ , for every  $0 \leq t \leq T-1$ ,

$$\left\| \frac{1}{r} \sum_{i=1}^r p_{t,i} p_{t,i}^\top \nabla f(x_t) \right\|^2 \leq \frac{Cd}{r} \log(cTr/\delta) \cdot \frac{1}{r} \sum_{i=1}^r |p_{t,i}^\top \nabla f(x_t)|^2$$

Furthermore,

$$\mathbb{P} \left( \bigcap_{t=0}^{T-1} \mathfrak{D}_t \right) \geq 1 - \delta$$

*Proof.* Write  $P = \begin{bmatrix} p_{t,1} & \cdots & p_{t,r} \end{bmatrix} \in \mathbb{R}^{d \times r}$ . By choosing  $s = \sqrt{1/c \log(2/\delta)}$  above, we know that

$$\frac{1}{\sqrt{r}} \|P\|_{\text{op}} \leq 1 + \sqrt{\frac{d}{r}} + \frac{1}{c\sqrt{r}} \sqrt{\log(c/\delta)} \leq C \sqrt{\frac{d}{r}} \sqrt{\log(c/\delta)}$$

with probability at least  $1 - \delta$ . As  $PP^\top = \sum_{i=1}^r p_{t,i} p_{t,i}^\top$ , we see that

$$\left\| \frac{1}{r} PP^\top \nabla f(x_t) \right\|^2 \leq \left( \frac{1}{\sqrt{r}} \|P\|_{\text{op}} \right)^2 \frac{1}{r} \left\| P^\top \nabla f(x_t) \right\|^2 \leq \frac{Cd}{r} \log(c/\delta) \frac{1}{r} \sum_{i=1}^r |p_{t,i}^\top \nabla f(x_t)|^2$$

We then rescale  $\delta$  to  $\delta/Tr$  to complete the proof.  $\square$

**Lemma 4.7.** *Define the big progress event*

$$\mathfrak{P}_t = \left\{ \frac{1}{r} \sum_{i=1}^r |\langle p_{t,i}, \nabla f(x_t) \rangle|^2 \geq \frac{1}{2} \|\nabla f(x_t)\|^2 \right\}$$

As long as  $r \geq 16\chi$ , we have

$$\mathbb{P} \left( \bigcap \mathfrak{P}_t \right) \geq 1 - \delta.$$

*Proof.* As  $\langle p_{t,i}, \nabla f(x_t) \rangle \stackrel{i.i.d.}{\sim} \mathcal{N}(0, \|\nabla f(x_t)\|^2)$ , by the Laurent-Massart inequality (Laurent and Massart, 2000, Lemma 1, Equation (4.4)),

$$\mathbb{P} \left( \chi_r^2 \leq \frac{r}{2} \right) \leq \exp \left( -\frac{r}{16} \right)$$

Therefore as long as  $r \geq 16\chi \geq 16 \log(T/\delta)$ ,

$$\frac{1}{r} \sum_{i=1}^r |\langle p_{t,i}, \nabla f(x_t) \rangle|^2 \geq \frac{1}{2} \|\nabla f(x_t)\|^2$$

for all  $0 \leq t \leq T-1$ . □

**Discussion.** It is unfortunate to necessitate a lower bound on  $r$  which will eventually grow with  $\varepsilon$ . The way around this is usually to instead look at the gradient progress over a time interval  $I$  of length  $\asymp \log(T/\delta)$ , in which case we could apply the same bound as above. However, this would require the gradients to be of roughly the same size, otherwise one gradient could dominate and we would only really have  $r$  effective samples even over the whole interval (compared to  $r \log(T/\delta)$ ).

When proving gradient descent converges to a first-order stationary point (in the sequel), this is possible, and alleviates this requirement on  $r$ . However, in the mixing case, since the  $z_t$  point depends only on  $z_{t-1}$  and not on  $y_{t-1}$ , the mirrored sequence could drift extremely far from the gradient sequence. In this case,  $x_{t+1} - x_t = (1 - \lambda)(y_t - x_t) + \lambda z_t - \lambda x_t$  could potentially grow unbounded. For example, consider the one-variable function  $f(x) = Gx + \frac{L}{2}x^2$ . Assume that  $x_t = 0$ , we know that  $\nabla f(x_t) = G$ . Then  $y_t = -\eta G$ .  $z_t = z_{t-1} - \alpha G$ . Then

$$|x_{t+1} - x_t| = |x_{t+1}| = |(1 - \lambda) - \eta G + \lambda(z_{t-1} - \alpha G)|$$

Sending  $|z_{t-1}| \rightarrow \infty$ , this gap can become unbounded. The  $\alpha$  dictating  $z_{t-1}$ 's progression grows with  $\varepsilon$ , so this is (a priori) not an unrealistic scenario. Furthermore, we will restart the algorithm later, with  $z_0 \neq x_0$ , so  $z_0$  really could be far away from  $y_0$ .

**Lemma 4.8.** *Define the bounded-perturbation event*

$$\mathfrak{B}_{t+1} := \bigcap_{i=1}^r \left\{ \sqrt{dL_{\delta/r}T} \leq \|p_{t+1,i}\| \leq \sqrt{3d\chi} \right\}$$

Then on  $\mathfrak{B}_{t+1} \cap \mathfrak{P}_{t+1}$ , as long as  $\eta \leq r/(12Ld\chi)$ , we have

$$f(y_{t+1}) \leq f(x_{t+1}) - \frac{\eta}{8} \|\nabla f(x_{t+1})\|^2 + C' L^2 \sigma^2 \chi^3 d^3$$

Furthermore,

$$\mathbb{P}\left(\bigcap_{t=0}^{T-1} \mathfrak{B}_t\right) \geq 1 - \delta$$

*Proof.* By  $L$ -smoothness,

$$\begin{aligned} f(y_{t+1}) &\leq f(x_{t+1}) + \langle \nabla f(x_{t+1}), y_{t+1} - x_{t+1} \rangle + \frac{L}{2} \|y_{t+1} - x_{t+1}\|^2 \\ &= f(x_{t+1}) - \eta \langle \nabla f(x_{t+1}), \tilde{\nabla} f(x_{t+1}) \rangle + \frac{L\eta^2}{2} \|\tilde{\nabla} f(x_{t+1})\|^2 \end{aligned}$$

Writing  $\tilde{\nabla} f(x_{t+1}) = \frac{1}{r} \sum_{i=1}^r p_i p_i^\top \nabla f(x_{t+1}) + R_{t+1}$ , we get

$$\begin{aligned} f(y_{t+1}) &\leq f(x_{t+1}) - \frac{\eta}{r} \sum_{i=1}^r |p_i^\top \nabla f(x_{t+1})|^2 - \eta \langle \nabla f(x_{t+1}), R_{t+1} \rangle \\ &\quad + L\eta^2 \left\| \frac{1}{r} \sum_{i=1}^r p_i p_i^\top \nabla f(x_{t+1}) \right\|^2 + L\eta^2 \|R_{t+1}\|^2 \end{aligned}$$

Now,  $2\langle \eta \nabla f(x_{t+1}), R_{t+1} \rangle \leq \eta^2 \|\nabla f(x_{t+1})\|^2 + R_{t+1}^2$  by Cauchy-Schwarz and Young's inequality.

Define the event where  $p_i$  are bounded:

$$\mathfrak{B}_{t+1} := \bigcap_{i=1}^r \left\{ \sqrt{dL\delta/rT} \leq \|p_{t+1,i}\| \leq \sqrt{3d\chi} \right\}$$

We know that  $\mathbb{P}(\mathfrak{B}_{t+1}) \geq 1 - \delta/rT$  by Lemma 5.7. Thus,

$$\begin{aligned} f(y_{t+1}) &\leq f(x_{t+1}) - \frac{\eta}{r} \sum_{i=1}^r |p_i^\top \nabla f(x_{t+1})|^2 + \eta^2 \gamma \|\nabla f(x_{t+1})\|^2 + (L\eta^2 + \gamma) \|R_{t+1}\|^2 \\ &\quad + \frac{CLd\eta^2 \log(Tr/\delta)}{r} \cdot \frac{1}{r} \sum_{i=1}^r |p_i^\top \nabla f(x_{t+1})|^2 \end{aligned}$$

From quadratic error lemma we know that  $|R_{t+1}| \leq 3L\sigma \log^{3/2}(1/\delta)d^{3/2}$ . Now by choosing  $\eta \leq r/(4CLd\chi)$ , we know that  $-\eta + CLd\chi\eta^2/r \leq -3/4\eta$ . As  $L\eta^2 < 1$  as  $L \geq 1$ , possibly

changing the absolute constant  $C'$ , we get,

$$f(y_{t+1}) \leq f(x_{t+1}) - \frac{3\eta}{4} \frac{1}{r} \sum_{i=1}^r |p_i^\top \nabla f(x_{t+1})|^2 + \eta^2 \|\nabla f(x_{t+1})\|^2 + C' L^2 \sigma^2 \chi^3 d^3$$

As long as  $\eta \leq 1/8$ ,  $-3\eta/8 + \eta^2 \leq -\eta/4$ , so using the previous lemma on  $\mathfrak{P}_{t+1}$ , we ensure

$$f(y_{t+1}) \leq f(x_{t+1}) - \frac{\eta}{8} \|\nabla f(x_{t+1})\|^2 + C' L^2 \sigma^2 \chi^3 d^3$$

□

#### 4.4 Mirror Descent against Progress

We can copy the mirror descent lemma from last time:

**Lemma 4.9.** *If  $z_{t+1} = \text{Mirr}_{z_t}(\alpha \tilde{\nabla} f(x_{t+1}))$ , then*

$$\forall u \in \mathbb{R}^d \quad \alpha \langle \tilde{\nabla} f(x_{t+1}), z_t - u \rangle \leq \frac{\alpha^2}{2} \left\| \tilde{\nabla} f(x_{t+1}) \right\|^2 + V_{z_t}(u) - V_{z_{t+1}}(u).$$

The standard coupling lemma comes from using  $\|\nabla f(x_{t+1})\|^2 \leq 2L(f(x_{t+1}) - f(y_{t+1}))$ . However this doesn't work here, as we are using an estimator. But we can fix that using the following lemma.

**Lemma 4.10.** *Define the event*

$$\mathfrak{E}_{t+1} = \left\{ \left\| \frac{1}{r} \sum_{i=1}^r p_{t+1,i} p_{t+1,i}^\top \nabla f(x_{t+1}) \right\|^2 \leq \frac{Cd}{r} \chi \|\nabla f(x_{t+1})\|^2 \right\}$$

On  $\mathfrak{B}_{t+1} \cap \mathfrak{E}_{t+1} \cap \mathfrak{D}_{t+1} \cap \mathfrak{P}_{t+1}$ , we have for all slack  $\gamma \geq 2$ ,

$$\frac{Cd\gamma\chi}{r} \|\nabla f(x_{t+1})\|^2 + \left\| \tilde{\nabla} f(x_{t+1}) \right\|^2 \leq \frac{16Cd\gamma\chi}{r\eta} (f(x_{t+1}) - f(y_{t+1}) + C' L^2 \sigma^2 \chi^3 d^3)$$

Furthermore,

$$\mathbb{P}\left(\bigcap \mathfrak{E}_t\right) \geq 1 - \delta.$$

*Proof.* Define the event

$$\mathfrak{E}_{t+1} = \left\{ \left\| \frac{1}{r} \sum_{i=1}^r p_{t+1,i} p_{t+1,i}^\top \nabla f(x_{t+1}) \right\|^2 \leq \frac{Cd}{r} \chi \|\nabla f(x_{t+1})\|^2 \right\}$$

Then on  $\mathfrak{E}_{t+1}$ ,

$$\begin{aligned} \|\tilde{\nabla} f(x_{t+1})\|^2 &\leq 2 \left\| \frac{1}{r} \sum_{i=1}^r p_i p_i^\top \nabla f(x_{t+1}) \right\|^2 + 2 \|R_{t+1}\|^2 \\ &\leq 2 \frac{Cd}{r} \chi \cdot \|\nabla f(x_{t+1})\|^2 + 2 \|R_{t+1}\|^2 \\ &\leq \frac{2Cd}{r} \chi \|\nabla f(x_{t+1})\|^2 + C' L^2 \sigma^2 \log^3(rT/\delta) d^3 \end{aligned}$$

We will need some slack of the gradient norm to tune later. For  $\gamma \geq 2$ , as  $\gamma - \gamma/2 \geq 1$ , this is at most

$$\frac{2Cd\gamma}{r} \chi \|\nabla f(x_{t+1})\|^2 + C' L^2 \sigma^2 \log^3(rT/\delta) d^3 - \frac{Cd\gamma}{r} \chi \|\nabla f(x_{t+1})\|^2$$

From before we know that

$$\begin{aligned} &\frac{\eta}{8} \|\nabla f(x_{t+1})\|^2 \\ &\leq f(x_{t+1}) - f(y_{t+1}) + C' L^2 \sigma^2 \log^3(rT/\delta) d^3 \end{aligned}$$

Therefore an upper bound on  $\|\tilde{\nabla} f(x_{t+1})\|^2$  is

$$\begin{aligned} &\frac{16Cd\gamma\chi}{r\eta} [f(x_{t+1}) - f(y_{t+1}) + C' L^2 \sigma^2 \log^3(rT/\delta) d^3] \\ &+ C' L^2 \sigma^2 \chi^3 d^3 - \frac{Cd\gamma}{r} \chi \|\nabla f(x_{t+1})\|^2 \end{aligned}$$

As  $\eta < 1/8$  we can absorb the last term into the first by changing the constant  $C'$ . Finally, after an application of the triangle inequality and Lemma 3.10, we have  $\mathbb{P}(\mathfrak{E}_{t+1}) \geq 1 - \delta/T$ , which gives the statement of the lemma.  $\square$



## 4.5 Linear Mixing

This lemma explains how to choose the mixing parameter  $\lambda$ .

**Lemma 4.11.** *When  $\lambda$  is chosen so  $\frac{1-\lambda}{\lambda} = \frac{8Cd\gamma\chi\alpha}{\eta}$ , the “regret” is bounded by*

$$\alpha\langle\nabla f(x_{t+1}), x_{t+1} - u\rangle - \alpha\langle\nabla f(x_{t+1}), z_t - u\rangle \leq \frac{8Cd\gamma\chi\alpha^2}{\eta}(f(y_t) - f(x_{t+1}))$$

*Proof.* We have by convexity and the definition of  $x_{t+1} = \lambda z_t + (1 - \lambda)y_t$ ,

$$\begin{aligned} & \alpha\langle\nabla f(x_{t+1}), x_{t+1} - u\rangle - \alpha\langle\nabla f(x_{t+1}), z_t - u\rangle \\ &= \alpha\langle\nabla f(x_{t+1}), x_{t+1} - z_t\rangle = \frac{(1 - \lambda)\alpha}{\lambda}\langle\nabla f(x_{t+1}), y_t - x_{t+1}\rangle \\ &\leq \frac{(1 - \lambda)\alpha}{\lambda}(f(y_t) - f(x_{t+1})) \end{aligned}$$

Now let  $\lambda \in (0, 1)$  satisfy  $\frac{1-\lambda}{\lambda} = \frac{8Cd\gamma\chi\alpha}{\eta}$  and we are done.  $\square$

This lemma reveals the martingale term we need to bound.

**Lemma 4.12.**

$$\begin{aligned} \alpha\langle\nabla f(x_{t+1}), z_t - u\rangle &\leq \alpha\langle\tilde{\nabla} f(x_{t+1}), z_t - u\rangle + \alpha\sum_{i=1}^r\langle(I - p_i p_i^\top)\nabla f(x_{t+1}), z_t - u\rangle \\ &\quad + C'T\alpha^2\|R_{t+1}\|^2 + \frac{1}{4T}V_{z_t}(u). \end{aligned}$$

*Proof.*

$$\langle\nabla f(x_{t+1}), z_t - u\rangle = \langle\tilde{\nabla} f(x_{t+1}), z_t - u\rangle + \sum_{i=1}^r\langle(I - p_i p_i^\top)\nabla f(x_{t+1}), z_t - u\rangle + \langle R_{t+1}, z_t - u\rangle$$

Now, by Young’s inequality,  $\alpha\langle R_{t+1}, z_t - u\rangle \leq C'\alpha^2 T\|R_{t+1}\|^2 + \frac{1}{8T}\|z_t - u\|^2$ , by strong convexity of Bregman divergence,  $V_{z_t}(u) \geq \frac{1}{2}\|z_t - u\|^2$ , so we arrive at the statement of the lemma.  $\square$

## 4.6 Dealing with the Martingale

The main idea from here is that, the martingale term

$$\sum_{t=0}^{T-1} \sum_{i=1}^r \langle (I - p_i p_i^\top) \nabla f(x_{t+1}), z_t - u \rangle$$

can be bounded by the “Energy”  $V_{z_t}(x^*)$ , which itself can be bounded by the sum of the gradient norms  $\sum \|\nabla f(x_t)\|^2$  and  $\Theta \geq V_{z_0}(x^*)$ .

This highlights a very general principle. To deal with martingale error terms with high probability, try bounding them by some other part of the equation which should have typical size  $\Theta(1)$ . In this case, that will be  $V_{z_t}(x^*)$  for us.

**Lemma 4.13.** *Define*

$$\mathfrak{M}_{t+1} = \left\{ \left| \sum_{t=0}^{\tau-1} \sum_{i=1}^r \alpha \langle (I - p_i p_i^\top) \nabla f(x_{t+1}), z_t - u \rangle \right| \leq \alpha \sqrt{Cr \sum_{t=0}^{\tau-1} \chi^2 \|\nabla f(x_t)\|^2 V_{z_t}(u)} \right\}$$

*Then*

$$\mathbb{P} \left( \bigcap_{t=0}^{T-1} \mathfrak{M}_t \right) \geq 1 - \delta.$$

*Proof.* From Lemma 3.10 above, we know that

$$\begin{aligned} & \left| \sum_{t=0}^{\tau-1} \sum_{i=1}^r \alpha \langle (I - p_i p_i^\top) \nabla f(x_{t+1}), z_t - u \rangle \right| \\ & \leq C\theta \sum_{t=0}^{\tau-1} \sum_{i=1}^r \log(CrT/\delta) \|\nabla f(x_{t+1})\|^2 \|z_t - u\|^2 + \frac{1}{\theta} \log\left(\frac{CrT}{\delta}\right) \end{aligned}$$

Using  $V_{z_t}(u) \geq \frac{1}{2} \|z_t - u\|^2$  and balancing the two terms by carefully choosing  $\theta$  according to AM-GM  $2\sqrt{ab} \leq a + b$ , we arrive at the statement of the lemma.  $\square$

## 4.7 Putting it all Together

First, as in the proof of Mirror Descent, we bound the regret, then  $V_{z_t}(x^*)$ , then last we will tune  $\alpha$ .

**Lemma 4.14** (Regret Bound). *On  $\bigcap_{t=1}^T (\mathfrak{B}_t \cap \mathfrak{D}_t \cap \mathfrak{P}_t \cap \mathfrak{G}_t \cap \mathfrak{M}_t)$ , we have*

$$\begin{aligned}
& \sum_{t=0}^{\tau-1} \alpha \langle \nabla f(x_t), x_t - x^* \rangle \\
& \leq \frac{8Cd\gamma\chi}{r\eta} \alpha^2 (f(y_0) - f(y_\tau)) - \frac{C\alpha^2 d\gamma\chi}{2r} \sum_{t=0}^{\tau-1} \|\nabla f(x_{t+1})\|^2 \\
& + \alpha \sqrt{\frac{C}{r} \sum_{t=0}^{\tau-1} \chi^2 \|\nabla f(x_{t+1})\|^2 V_{z_t}(x^*)} + \frac{C'(\alpha^2 T + \alpha T^2) \gamma L^2 \sigma^2 \chi^4 d^4}{r\eta} \\
& + V_{z_0}(x^*) - V_{z_\tau}(x^*) + \frac{1}{4T} \sum_{t=0}^{\tau-1} V_{z_t}(x^*)
\end{aligned}$$

*Proof.* By Lemma 4.11, we know that

$$\sum_{t=0}^{\tau-1} \alpha \langle \nabla f(x_t), x_t - x^* \rangle \leq \sum_{t=0}^{\tau-1} \frac{8Cd\gamma\chi\alpha^2}{\eta} (f(y_t) - f(x_{t+1})) + \sum_{t=0}^{\tau-1} \alpha \langle \nabla f(x_{t+1}), z_t - u \rangle$$

By Lemma 4.12, we know that

$$\begin{aligned}
& \sum_{t=0}^{\tau-1} \alpha \langle \nabla f(x_{t+1}), z_t - u \rangle \\
& \leq \alpha \sum_{t=0}^{\tau-1} \left[ \langle \tilde{\nabla} f(x_{t+1}), z_t - u \rangle + \frac{1}{r} \sum_{i=1}^r \langle (I - p_i p_i^\top) \nabla f(x_{t+1}), z_t - u \rangle + 2T \|R_{t+1}\|^2 + \frac{1}{T} V_{z_t}(u) \right]
\end{aligned}$$

From above we know that

$$\alpha \sum_{t=0}^{\tau-1} \sum_{i=1}^r \langle (I - p_i p_i^\top) \nabla f(x_{t+1}), z_t - u \rangle \leq \alpha \sqrt{Cr \sum_{t=0}^{\tau-1} \chi^2 \|\nabla f(x_{t+1})\|^2 V_{z_t}(u)}$$

From Lemmas 4.9 and 4.10 we know that

$$\begin{aligned}
& \sum_{t=0}^{\tau-1} \alpha \langle \tilde{\nabla} f(x_{t+1}), z_t - u \rangle \\
& \leq \alpha^2 \sum_{t=0}^{\tau-1} \frac{8Cd\gamma\chi}{r\eta} (f(x_{t+1}) - f(y_{t+1}) + C'L^2\sigma^2\chi^3d^3) - \frac{C\alpha^2d\gamma\chi}{2r} \sum_{t=0}^{\tau-1} \|\nabla f(x_{t+1})\|^2 \\
& + V_{z_0}(x^*) - V_{z_\tau}(x^*)
\end{aligned}$$

Recall that  $\|R_{t+1}\|^2 \leq C'L^2\sigma^2\chi^3d^3$ . Putting all of it together, we get:

$$\begin{aligned}
& \sum_{t=0}^{\tau-1} \alpha \langle \nabla f(x_t), x_t - x^* \rangle \\
& \leq \frac{8Cd\gamma\chi}{r\eta} \alpha^2 (f(y_0) - f(y_\tau)) - \frac{C\alpha^2d\gamma\chi}{2r} \sum_{t=0}^{\tau-1} \|\nabla f(x_{t+1})\|^2 \\
& + \alpha \sqrt{\frac{C}{r} \sum_{t=0}^{\tau-1} \chi^2 \|\nabla f(x_{t+1})\|^2 V_{z_t}(x^*)} + \frac{C'\alpha^2T\gamma L^2\sigma^2\chi^4d^4}{r\eta} + C'\alpha T^2L^2\sigma^2\chi^3d^3 \\
& + V_{z_0}(x^*) - V_{z_\tau}(x^*) + \frac{1}{4T} \sum_{t=0}^{\tau-1} V_{z_t}(x^*)
\end{aligned}$$

By using  $r \leq d$  and  $\eta < 1$  to handle the two quadratic terms we arrive at the statement of the lemma.  $\square$

This next lemma is what we meant by bounding the martingale by other already-present  $\Theta(1)$  terms.

**Lemma 4.15** (Energy Bound). *On  $\bigcap_{t=1}^T (\mathfrak{B}_t \cap \mathfrak{D}_t \cap \mathfrak{P}_t \cap \mathfrak{G}_t \cap \mathfrak{M}_t)$ , upon defining  $M = \sup_{0 \leq t \leq T-1} V_{z_t}(x^*)$ , choosing the slack  $\gamma = 8\chi$  and letting  $\Theta \geq V_{z_0}(x^*)$  be any upper bound, we have*

$$\frac{M}{2} \leq \frac{256Cd\chi^2}{r\eta} \alpha^2 (f(y_0) - f^*) + \Theta + \frac{C'(\alpha^2T + \alpha T^2)L^2\sigma^2\chi^4d^4}{r\eta}$$

*Proof.* By convexity we know that  $f^* = f(x^*) \geq f(x) + \langle \nabla f(x), x^* - x \rangle$ . Rearranging, we

get  $\langle \nabla f(x), x^* - x \rangle \geq f(x) - f^* \geq 0$ . Using Young's inequality, we know that

$$\alpha \sqrt{\frac{C}{r} \chi^2 \sum_{t=0}^{\tau-1} \|\nabla f(x_{t+1})\|^2} V_{z_t}(x^*) \leq \frac{M}{4} + \alpha^2 \frac{4C\chi^2}{r} \sum_{t=0}^{\tau-1} \|\nabla f(x_{t+1})\|^2$$

Therefore by the previous lemma, and using  $\frac{1}{T} \sum_{t=0}^{\tau-1} V_{z_t}(x^*) \leq \frac{1}{4}M$ , we know that

$$\begin{aligned} V_{z_\tau}(x^*) &\leq \frac{8C\gamma d\chi}{r\eta} \alpha^2 (f(y_0) - f(y_\tau)) - \frac{C\alpha^2 d\gamma\chi}{2r} \sum_{t=0}^{\tau-1} \|\nabla f(x_{t+1})\|^2 \\ &\quad + V_{z_0}(x^*) + \frac{M}{4} + \frac{C'(\alpha^2 T + \alpha T^2) \gamma L^2 \sigma^2 \chi^4 d^4}{r\eta} \\ &\quad + \alpha^2 \frac{4C\chi}{r} \sum_{t=0}^{\tau-1} \|\nabla f(x_{t+1})\|^2 + \frac{M}{4} \end{aligned}$$

By choosing  $\gamma = 8\chi$ , as  $d \geq 1$ , we know that

$$-\frac{C\alpha^2 d\gamma\chi}{2r} \sum_{t=0}^{\tau-1} \|\nabla f(x_{t+1})\|^2 + \alpha^2 \frac{4C\chi^2}{r} \sum_{t=0}^{\tau-1} \|\nabla f(x_{t+1})\|^2 \leq 0$$

By using  $f(y_0) - f(y_\tau) \leq f(y_0) - f^*$  along with  $V_{z_0}(x^*) \leq \Theta$ , and taking supremum of both sides, we have

$$M \leq \frac{64Cd\chi^2}{r\eta} \alpha^2 (f(y_0) - f^*) + \Theta + \frac{M}{2} + \frac{C'(\alpha^2 T + \alpha T^2) L^2 \sigma^2 \chi^4 d^4}{r\eta}$$

which completes the proof.  $\square$

**Theorem 4.16** (Halving Function Objective). *Assuming the initial error  $f(x_0) - f^* \leq \Delta$ , and defining  $\bar{x} = 1/T \sum_{t=0}^{T-1} x_t$ , by choosing*

$$\eta = \frac{r}{4CLd\chi}, \quad \alpha = \frac{r}{16Cd\chi^{3/2}} \sqrt{\frac{\Theta}{L\Delta}}, \quad \sigma^2 \leq \frac{r^3 \Delta^{3/2}}{4C' L^{7/2} \chi^{15/2} d^6 \sqrt{\Theta}}, \quad r \geq 16\chi$$

In

$$T = \frac{256Cd\chi^{3/2}}{r} \sqrt{\frac{L\Theta}{\Delta}}$$

iterations we have  $f(\bar{x}) - f^* \leq \Delta/2$  with probability at least  $1 - 5\delta$ .

*Proof.* On  $\bigcap_{t=1}^T (\mathfrak{B}_t \cap \mathfrak{D}_t \cap \mathfrak{P}_t \cap \mathfrak{G}_t \cap \mathfrak{M}_t)$  all our previous lemmas hold. Each has probability

at least  $1 - \delta$  so the intersection of 5 has probability at least  $1 - 5\delta$ .

By Lemma 4.14, we know that

$$\begin{aligned}
& \sum_{t=0}^{T-1} \alpha \langle \nabla f(x_t), x_t - x^* \rangle \\
& \leq \frac{64Cd\chi}{r\eta} \alpha^2 (f(y_0) - f(y_T)) - \frac{C\alpha^2 d\gamma\chi}{2r} \sum_{t=0}^{T-1} \|\nabla f(x_{t+1})\|^2 \\
& + \alpha \sqrt{\frac{C}{r} \sum_{t=0}^{T-1} \chi^2 \|\nabla f(x_{t+1})\|^2} V_{z_t}(x^*) + \frac{C'(\alpha^2 T + \alpha T^2) L^2 \sigma^2 \chi^4 d^4}{r\eta} \\
& + V_{z_0}(x^*) - V_{z_T}(x^*) + \frac{1}{4T} \sum_{t=0}^{T-1} V_{z_t}(x^*)
\end{aligned}$$

Using Young's inequality in exactly the way as the previous lemma, we can rearrange this to

$$\begin{aligned}
& \sum_{t=0}^{T-1} \alpha \langle \nabla f(x_t), x_t - x^* \rangle \\
& \leq \frac{64Cd\chi^2}{r\eta} \alpha^2 (f(y_0) - f^*) + \frac{C'(\alpha^2 T + \alpha T^2) L^2 \sigma^2 \chi^4 d^4}{r\eta} + \Theta + M/4
\end{aligned}$$

Using the previous lemma, we have

$$\begin{aligned}
& \sum_{t=0}^{T-1} \alpha \langle \nabla f(x_t), x_t - x^* \rangle \\
& \leq \frac{128Cd\chi^2}{r\eta} \alpha^2 (f(y_0) - f^*) + \frac{C'(\alpha^2 T + \alpha T^2) L^2 \sigma^2 \chi^4 d^4}{r\eta} + 2\Theta
\end{aligned}$$

Let  $\bar{x} = \sum_{t=0}^{T-1} x_t$ . By convexity,

$$T(f(\bar{x}) - f^*) \leq \sum_{t=0}^{T-1} (f(x_t) - f^*) \leq \sum_{t=0}^{T-1} \langle \nabla f(x_t), x_t - x^* \rangle$$

Therefore, assuming the initial objective gap  $f(y_0) - f^* \leq \Delta$ ,

$$f(\bar{x}) - f^* \leq \frac{128Cd\chi^2 \alpha \Delta}{r\eta T} + \frac{2\Theta}{\alpha T} + \frac{C'(\alpha + T) L^2 \sigma^2 \chi^4 d^4}{r\eta}$$

Now choose the learning rate  $\eta = r/(4CLd\chi)$ . Then we get

$$f(\bar{x}) - f^* \leq \frac{512C^2d^2L\chi^3\alpha\Delta}{r^2T} + \frac{2\Theta}{\alpha T} + \frac{C'(\alpha + T)L^3\sigma^2\chi^5d^5}{r^2}$$

By choosing

$$\alpha = \frac{r}{16Cd\chi^{3/2}} \sqrt{\frac{\Theta}{L\Delta}},$$

and

$$T = \frac{256Cd\chi^{3/2}}{r} \sqrt{\frac{L\Theta}{\Delta}}$$

as  $L \geq 1$  we have  $T \geq \alpha$ , so the above is at most

$$\frac{\Delta}{4} + \sqrt{\frac{L\Theta}{\Delta}} \frac{C'L^3\sigma^2\chi^{15/2}d^6}{r^3}$$

Then as long as

$$\sigma^2 \leq \frac{r^3\Delta^{3/2}}{4C'L^{7/2}\chi^{15/2}d^6\sqrt{\Theta}}$$

We have

$$f(\bar{x}) - f^* \leq \frac{\Delta}{2}$$

thereby halving the objective value. This completes the proof.  $\square$

We want to restart this halving procedure  $\log(1/\varepsilon)$  times so we can converge to a point with objective gap  $f(\bar{x}) - f^* \leq \varepsilon$ . However, the  $\bar{x}$  returned might not satisfy  $V_{\bar{x}}(x^*) \leq \Theta$ . Indeed, this valuation could actually go up if there is multiple minima. This point is actually unmentioned and unresolved in Allen-Zhu and Orecchia (2016). The simple fix is to notice above that the requirement  $z_0 = y_0$  is not necessary. Actually, the only fact we used about  $z_0$  was that  $V_{z_0}(x^*) \leq \Theta$ . So by restarting the algorithm  $n = \log_2(\Delta/\varepsilon)$  times with  $x_0 = \bar{x} = y_0$ , and  $z_0$  as it's original value each time, we can find a point with

$f(\bar{x}) - f^* \leq \varepsilon$  in

$$T = \tilde{O} \left( \frac{d}{r} \sqrt{\frac{L\Theta}{\varepsilon}} + \frac{d}{r} \sqrt{\frac{L\Theta}{2\varepsilon}} + \frac{d}{r} \sqrt{\frac{L\Theta}{4\varepsilon}} + \dots \right) = \tilde{O} \left( \frac{d}{r} \sqrt{\frac{L\Theta}{\Delta}} \right)$$

iterations, which yields the following theorem:

**Theorem 4.17.** *By restarting Algorithm 2  $\log_2(\Delta/\varepsilon)$  many times, we find a point  $\bar{x}$  so that  $f(\bar{x}) - f^* \leq \varepsilon$  in*

$$T = \tilde{O} \left( \frac{d}{r} \sqrt{\frac{L\Theta}{\varepsilon}} \right)$$

*iterations.*

When  $r = d$ , this matches the asymptotically optimal running time to converge to an  $\varepsilon$ -optimal point for a  $L$ -smooth convex function, up to logarithmic factors. To the best of our knowledge, this is the first zero-order accelerated gradient descent algorithm for smooth convex functions that achieves the optimal running time (up to log factors) which can be run once and output a solution with high probability. On the contrary, expectation-based methods need to run  $\log(1/\delta)$  independent copies to ensure they work with high probability. We believe this analysis gives a strong intuition about the fluctuations one often sees when running the algorithm in practice, as we have directly compared the martingale term to the progress and distance to optimum.



## Chapter 5

# EGGROLL for Nonconvex Optimization

In previous chapters we have largely focused on the theoretical aspects of zero-order optimization methods. However, their implementation and practical use is also an important question. While we have focused on zero-order methods, how do large GPU clusters actually share the perturbations with each other? *A priori*, they would have to send the full weights of each perturbation to every cluster, which could take prohibitively long due to synchronization requirements. This would effectively eliminate the gains behind evaluating  $r$  perturbations in parallel we have discussed in detail above.

Salimans et al. (2017) solved this problem by sharing only the pseudorandom seeds used to generate each Gaussian perturbation. Full high-dimensional perturbations can thus be shared between GPUs using only a 32-bit seed, where the receiving GPU subsequently reconstructs the perturbation locally. For the first time, zeroth-order methods, also called Evolutionary Strategies (ES), was scaled to machine learning models with millions of parameters, using more than one thousand parallel workers. But as we have seen in

Chapter 4, increasing the population size  $r$  linearly speeds up the algorithm. A two-times increase in the number of parallel workers yields precisely a two-times speedup in training speed.

The bottleneck now, having solved the problem of sharing the perturbations across GPUs, was memory. Each Gaussian perturbation has  $d$  nonzero entries almost surely, where  $d$  is the dimension of the model. Thus, for population size  $r$ , each GPU needs to store  $O(d \cdot r)$  parameters. This fundamentally limits how big  $r$  can be in practice.

Very recently, Sarkar et al. (2025) introduced EGGROLL to solve this bottleneck by replacing full Gaussian perturbations  $P \sim \mathcal{N}(0, I_{d \times m})$  with rank- $r$  perturbations  $A \in \mathbb{R}^{d \times k}$  and  $B \in \mathbb{R}^{m \times k}$  with  $E = \frac{1}{\sqrt{k}}AB^\top$ , critically using  $u(M + \sigma P)^\top = uM^\top + \frac{\sigma}{\sqrt{k}}(uB)A^\top$ . Using this low-rank trick, they were able to scale ES vastly further than done previously. On the theoretical side, they showed that, under some reasonable assumptions if  $\sigma = o(d^{-1/2})$  then as  $d \rightarrow \infty$ ,

$$\frac{1}{\sigma} \mathbb{E}_{p \sim \mathcal{N}(0, I_d)} [f(x + \sigma P)P] \rightarrow \nabla f(x)$$

Furthermore, they prove that as the rank of the perturbations  $k \rightarrow \infty$ ,

$$\frac{1}{\sigma} \mathbb{E}_{E \sim p(E)} [Ef(M + \sigma E)] \rightarrow \frac{1}{\sigma} \mathbb{E}_{p \sim \mathcal{N}(0, I_{d \times m})} [f(x + \sigma P)P]$$

this is because the distribution of  $E$  will approach  $\mathcal{N}(0, I_{d \times m})$  with increasing rank.

## 5.1 Main Result

The preceding result suggests that increasing the rank of the perturbations is beneficial, so the resulting distribution becomes Gaussian. The Gaussian distribution is rotationally invariant and the rank-1 update  $ab^\top$  is not, maybe we need this property to “locate” the gradient. However, this is not true, it turns out that the requirement for convergence is

the isotropic condition mentioned previously in the preliminaries. We prove that, despite only using rank-1 perturbations  $ab^\top$ , not only do we get convergence to a first-order stationary point, but that its speed also scales linearly with the population size. The practical takeaway is that when function evaluations are performed in parallel, twice as much compute converges twice as fast.

**Assumption 5.1** (Properties of  $f$ ). We suppose that  $f : \mathbb{R}^{d \times m} \rightarrow \mathbb{R}$  satisfies the following properties:

1.  $f$  is differentiable and lower bounded, i.e.  $f^* = \min_x f(x) > -\infty$
2.  $f$  is  $L$ -gradient Lipschitz, i.e.

$$\|\nabla f(X) - \nabla f(Y)\|_F \leq L \|X - Y\|_F \quad \forall X, Y \in \mathbb{R}^{d \times m}$$

**Definition 5.2.** A point  $X \in \mathbb{R}^{d \times m}$  is an  $\varepsilon$ -first order stationary point if

$$\|\nabla f(X)\|_F \leq \varepsilon$$

In this excerpt we work with symmetric EGGROLL, following the low-rank perturbation framework of Sarkar et al. (2025).

**Definition 5.3** (Symmetric EGGROLL Update). We define a rank  $r$  symmetric zeroth order estimator as follows:

$$g_\sigma^r(X) := \frac{1}{r} \sum_{i=1}^r \frac{f(X + \sigma P_i) - f(X - \sigma P_i)}{2\sigma} P_i$$

using  $P_i = a_i b_i^\top$  where  $a_i \stackrel{i.i.d.}{\sim} \mathcal{N}(0, I_d)$  and  $b_i \stackrel{i.i.d.}{\sim} \mathcal{N}(0, I_m)$ , and  $\sigma > 0$  is the smoothing radius.

---

**Algorithm 3** Rank- $r$  Symmetric EGGROLL

---

**Require:** Initial point  $X_0 \in \mathbb{R}^{d \times m}$ , horizon  $T$ , stepsize  $\eta > 0$ , smoothing radius  $\sigma > 0$ , rank  $r \in \mathbb{N}$   
**for**  $t = 0, \dots, T - 1$  **do**  
    Sample  $\{a_{t,i}\}_{i=1}^r \sim \mathcal{N}(0, I)$  and  $\{b_{t,i}\}_{i=1}^r \sim \mathcal{N}(0, I)$  to compute  $g_\sigma^r(X_t)$   
    Update  $X_{t+1} \leftarrow X_t - \eta g_\sigma^r(X_t)$   
**end for**  
**return**  $X_T$

---

We state an informal version of the main result.

**Theorem 5.4** (informal). *Consider running Algorithm 3. Suppose  $0 < \delta \leq 1/e$ . Suppose*

$$\sigma = \tilde{O}\left(\frac{\varepsilon}{dm}\right), \quad \eta = \tilde{O}\left(\frac{r}{dm}\right).$$

*Then in*

$$T = \tilde{\Omega}\left(\frac{dmL(f(X_0) - f^*)}{r\varepsilon^2}\right)$$

*iterations (with each iteration using  $2r$  function evaluations), with probability at least  $1 - \delta$ , at least  $3/4$  of the iterates are  $\varepsilon$ -first order stationary points.*

The overall proof strategy follows that of Ren et al. (2023). The main new ingredient is Theorem 5.8, which establishes the concentration bound required when full Gaussian perturbations are replaced by rank-one perturbations. The Gaussian argument does not extend to this setting because rank-one perturbations have fat tails, with entries that are sub-exponential rather than sub-gaussian.

## 5.2 Concentration Bounds

Here we prove the novel concentration bound about the EGGROLL error. Throughout this excerpt, all norms of matrices is the Frobenius norm

$$\|A\|_F = \left( \sum_{i,j} A_{i,j}^2 \right)^{1/2}.$$

Throughout this paper, we denote

$$\mathbb{E}_{\mathcal{F}_{t-1}}[X; A] := \mathbb{E}[X \mathbb{I}_A \mid \mathcal{F}_{t-1}]$$

We have seen subGaussian random variables in the preliminaries. Here, we will need a generalization, called norm-subGaussian.

**Definition 5.5.** The random variable  $X$  is norm-subGaussian with parameter  $\sigma$ , abbreviated  $\text{nSG}(\sigma)$ , if

$$\mathbb{E} \left[ \exp \left( \frac{\|X\|^2}{3\sigma^2} \right) \right] \leq 2$$

As an example,  $p \sim \mathcal{N}(0, I_d)$  is  $\text{nSG}(\sqrt{d})$ . The next lemma is basically a Hoeffding-like inequality for norm-subGaussian random variables.

**Lemma 5.6** (Jin et al. (2019)). *Let random vectors  $x_1, \dots, x_n \in \mathbb{R}^d$ , and corresponding filtrations  $\mathcal{F}_t = \sigma(X_1, \dots, X_t)$  satisfy that*

$$\mathbb{E}_{\mathcal{F}_{t-1}}[x_t] = 0, \quad \mathbb{P}_{\mathcal{F}_{t-1}}(\|x_t\| \geq s) \leq 2e^{-\frac{s^2}{3\sigma^2}} \quad \forall s \in \mathbb{R}$$

*Then there exists an absolute constant  $c$  such for any fixed  $\delta > 0, \theta > 0$ , with probability at least  $1 - \delta$ ,*

$$\left\| \sum_{i=1}^n X_i \right\| \leq c \cdot \theta \sum_{i=1}^n \sigma_i^2 + \frac{1}{\theta} \log \left( \frac{2d}{\delta} \right).$$

This will help us later.

**Lemma 5.7.** *Let  $X$  follow  $\chi_d^2$  distribution and  $0 < \delta < 1/e$ . Then there exists  $L_\delta$  such that for every  $d$ , with*

$$A := \{dL_\delta^2 \leq X \leq 3d \log(C/\delta)\}$$

*we have*

$$\mathbb{E}[(X - d); A] = 0, \quad \mathbb{P}(A) \geq 1 - \delta$$

*for some absolute constant  $C$ .*

*Proof.* Let  $f_d(x)$  be the  $\chi_d^2$  density,

$$f_d(x) = \frac{1}{2^{d/2} \Gamma(d/2)} x^{d/2-1} e^{-x/2}, \quad x \geq 0.$$

A simple calculation shows that

$$(x - d)f_d(x) = -2 \frac{d}{dx} (x f_d(x))$$

Define  $U = 3d \log(1/\delta)$  and define the set

$$A = \{L \leq X \leq U\}$$

We want

$$\mathbb{E}[(X - d); A] = \int_L^U (x - d)f_d(x)dx = -2L f_d(L) - 2U f_d(U) \stackrel{\text{want}}{=} 0.$$

Therefore, we need

$$\begin{aligned} L f_d(L) &= U f_d(U) \\ \iff L^d e^{-L} &= U^d e^{-U} \\ \iff U - L &= d \log(U/L) \end{aligned}$$

Now define  $\ell = L/d$  and  $u = U/d$ . Then we want to solve

$$\begin{aligned} du - d\ell &= d \log(u/\ell) \\ \iff \ell e^{-\ell} &= u e^{-u} \end{aligned}$$

The function  $g(x) = x e^{-x}$  is increasing on  $[0, 1)$  and decreasing on  $(1, \infty)$ . Therefore, as long as  $u > 1$ , which follows as  $\delta < 1/e$ , there exists a unique  $\ell < 1$  satisfying the above equation. Now notice, as  $e^{-\ell} \geq e^{-1}$ ,

$$\ell e^{-1} \leq \ell e^{-\ell} = 3 \log(1/\delta) \delta^3 \leq 3 e^{-1} \delta^2$$

Now we have that

$$\mathbb{P}(A^c) = \mathbb{P}(X \geq U) + \mathbb{P}(X \leq L)$$

Recall the Chernoff bound For  $0 < \alpha < 1$ ,

$$\mathbb{P}(X \leq \alpha d) \leq (\alpha e^{1-\alpha})^{d/2} \leq (e\alpha)^{d/2}$$

Then we have

$$\mathbb{P}(X \leq L) \leq \mathbb{P}(X \leq 3d\delta^2) \leq (3e\delta^2)^{d/2} = (\sqrt{3e}\delta)^d$$

When  $\sqrt{3e}\delta < 1$ , this is  $< \sqrt{3e}\delta$ , otherwise this bound holds vacuously. Therefore  $\mathbb{P}(X \leq L) \leq \sqrt{3e}\delta$ . On the other hand, using that for  $t < 1/2$ ,

$$\mathbb{E}\left[e^{t\|Z\|^2}\right] = (1 - 2t)^{-d/2},$$

along with Markov's inequality, we get

$$\mathbb{P}(X \geq 3d \log(1/\delta)) = \mathbb{P}(\exp(X/3) \geq \exp(d \log(1/\delta))) \leq 3^{d/2} \delta^d \leq \sqrt{3}\delta$$

And thus  $\mathbb{P}(A) \geq 1 - C\delta$  for the absolute constant  $C = \sqrt{3} + \sqrt{3e}$ .  $\square$

Here is the novel concentration bound. The proof is inspired by Proposition 2 of Ren

et al. (2023) but with many important differences.

**Theorem 5.8.** *Let  $\{\mathcal{F}_t\}_{t \geq -1}$  be a filtration, and let  $\{a_t\}_{t \geq 0}, \{b_t\}_{t \geq 0}$  be two sequences of independent random vectors such that  $a_t \sim \mathcal{N}(0, I_d)$ ,  $b_t \sim \mathcal{N}(0, I_m)$ ,  $a_t, b_t \in \mathcal{F}_t$  and  $a_t, b_t$  are independent of  $\mathcal{F}_{t-1}$ . Let  $\{V_t\}_{t \geq 0}$  be a sequence of random matrices such that  $V_t \in \mathcal{F}_{t-1}$ . For each  $\tau \geq 0$ , let*

$$E_\tau = \sum_{t=0}^{\tau-1} \left( V_t - \langle V_t, a_t b_t^\top \rangle a_t b_t^\top \right)$$

*Then, there exists some absolute constants  $C > 0, c > 0$  such that for any  $\tau \in \mathbb{Z}^+$ , and  $0 < \delta < 1/e$ , for any  $\theta > 0$ , with probability at least  $1 - \delta$  given  $\mathcal{F}_{t-1}$ ,*

$$\|E_\tau\|_F \leq c \cdot \theta \sum_{t=0}^{\tau-1} dm \log^3 \left( \frac{C\tau m}{\delta} \right) \|V_t\|_F^2 + \frac{1}{\theta} \log \left( \frac{Cm^2 d \tau}{\delta} \right).$$

*Proof.* Let  $V_t = P\Sigma Q^\top = \sum_{i=1}^r \sigma_i p_i q_i^\top$  be the singular value decomposition of  $V_t$ . Extend  $\{q_1, \dots, q_r\}$  to an orthonormal basis  $\{q_1, \dots, q_m\}$  of  $\mathbb{R}^m$ . Define the events

$$A_t = \left\{ \sqrt{d}L_\delta < \|a_t\| < \sqrt{3d \log(C/\delta)} \right\}$$

and

$$B_{t,i} = \left\{ L_\delta < |b_t^\top q_i| < \sqrt{3 \log(C/\delta)} \right\}, \quad B_t = \bigcap_{i=1}^m B_{t,i}.$$

Let  $R$  be an orthogonal matrix  $RR^\top = R^\top R = I$ , and notice that  $Ra$  has the same distribution as  $a$ , and that  $A_t$  is the same event when  $\|a_t\|$  is replaced with  $\|Ra_t\|$ . Therefore,

$$R \mathbb{E}_{\mathcal{F}_{t-1}} \left[ a_t a_t^\top; A_t \right] R^\top = \mathbb{E}_{\mathcal{F}_{t-1}} \left[ (Ra_t)(Ra_t)^\top; A_t \right] = \mathbb{E}_{\mathcal{F}_{t-1}} \left[ a_t a_t^\top; A_t \right]$$

Now, if  $M$  is a symmetric matrix such that  $RM R^\top = M$  for all  $R$ , it must be that  $M$  is a multiple of the identity. Therefore,

$$\mathbb{E}_{\mathcal{F}_{t-1}} \left[ a_t a_t^\top; A_t \right] = \alpha I_d.$$



We then get

$$\begin{aligned}\alpha d &= \text{Tr}\left(\mathbb{E}_{\mathcal{F}_{t-1}}\left[a_t a_t^\top; A_t\right]\right) = \mathbb{E}_{\mathcal{F}_{t-1}}\left[\|a_t\|^2; A_t\right] \\ &= \mathbb{E}_{\mathcal{F}_{t-1}}\left[(\|a_t\|^2 - d); A_t\right] + d\mathbb{P}(A_t) = d\mathbb{P}(A_t)\end{aligned}$$

The most important fact for this entire proof, is to notice that the events  $B_{t,i}$  are independent.

This is because, since  $q_i \perp q_j$  for  $i \neq j$ , the  $b_t^\top q_i$  and  $b_t^\top q_j$  are independent given  $\mathcal{F}_{t-1}$ .

Thus, if  $i \neq j$ ,

$$\begin{aligned}\mathbb{E}_{\mathcal{F}_{t-1}}\left[(Q^\top b_t)(Q^\top b_t)_{i,j}^\top; B_t\right] &= \mathbb{E}_{\mathcal{F}_{t-1}}\left[(q_i^\top b_t)(q_j^\top b_t); B_t\right] \\ &= \mathbb{P}_{\mathcal{F}_{t-1}}\left(\bigcap_{k \neq i,j} B_{t,k}\right) \cdot \mathbb{E}_{\mathcal{F}_{t-1}}\left[q_i^\top b_t; B_{t,i}\right] \cdot \mathbb{E}_{\mathcal{F}_{t-1}}\left[q_j^\top b_t; B_{t,j}\right] \\ &= 0\end{aligned}$$

because if  $Z$  is  $\mathcal{N}(0, 1)$  and  $A \in \sigma(Z)$  is a symmetric event, we have  $\mathbb{E}[Z; A] = 0$ . When  $i = j$ ,

$$\mathbb{E}_{\mathcal{F}_{t-1}}\left[(q_i^\top b_t)^2; B_t\right] = \mathbb{P}_{\mathcal{F}_{t-1}}\left(\bigcap_{k \neq i} B_{t,k}\right) \cdot \mathbb{E}_{\mathcal{F}_{t-1}}\left[(q_i^\top b_t)^2; B_{t,i}\right] = \mathbb{P}_{\mathcal{F}_{t-1}}(B_t)$$

by our judicious choice of  $B_{t,i}$  (cf. the lemma before this).

As  $Q \in \mathcal{F}_{t-1}$ , we thus have

$$\begin{aligned}\mathbb{E}_{\mathcal{F}_{t-1}}\left[Q^\top b_t b_t^\top; B_t\right] &= \mathbb{E}_{\mathcal{F}_{t-1}}\left[(Q^\top b_t)(Q^\top b_t)^\top; B_t\right] \cdot Q^\top \\ &= \mathbb{P}_{\mathcal{F}_{t-1}}(B)I \cdot Q^\top = \mathbb{P}_{\mathcal{F}_{t-1}}(B)Q^\top.\end{aligned}$$

Therefore,

$$\mathbb{E}_{\mathcal{F}_{t-1}}\left[V_t - a_t a_t^\top V_t b_t b_t^\top; A_t, B_t\right] = V_t \mathbb{P}_{\mathcal{F}_{t-1}}(A_t \cap B_t) - \mathbb{E}_{\mathcal{F}_{t-1}}\left[a_t a_t^\top P \Sigma Q^\top b_t b_t^\top; A_t, B_t\right]$$

Now,

$$\begin{aligned}\mathbb{E}_{\mathcal{F}_{t-1}}[a_t a_t^\top P \Sigma Q^\top b_t b_t^\top; A_t, B_t] &= \mathbb{E}_{\mathcal{F}_{t-1}}[a_t a_t^\top; A_t] P \Sigma \cdot \mathbb{E}_{\mathcal{F}_{t-1}}[Q^\top b_t b_t^\top; B_t] \\ &= P \Sigma Q^\top = V_t\end{aligned}$$

Therefore, upon defining

$$Q_t = (V_t - a_t a_t^\top V_t b_t b_t^\top) \mathbb{1}_{A_t \cap B_t}$$

we have

$$\mathbb{E}_{\mathcal{F}_{t-1}}[Q_t] = 0.$$

The next part of the proof is showing that  $Q_t$  is  $\text{nSG}(C\sqrt{mn}\text{polylog}(dm/\delta) \|V\|_F)$ .

First, notice that on  $B_t$ ,

$$\|b_t\|^2 = \sum_{i=1}^m |q_i^\top b_t| \leq 3m \log(C/\delta)$$

and recall that  $A_t$  was defined so that

$$\|a_t\| \leq \sqrt{3d \log(C/\delta)}.$$

So on  $A_t \cap B_t$  we have

$$\left\| a_t a_t^\top V_t b_t b_t^\top \right\|_F = |a_t^\top V b_t| \|a_t\| \|b_t\| \leq |a_t^\top V b_t| \sqrt{9md \log^2(C/\delta)}$$

Our main goal now is to show that  $a_t^\top V b_t$  is subgaussian with parameter  $\tilde{O}(\|V\|_F)$ . To simplify notation, we will drop the  $ts$  and conditioning, it won't be important for the argument. Recall that for  $Z \sim \mathcal{N}(0, \sigma^2)$ , if  $t\sigma^2 < \frac{1}{2}$ ,

$$\mathbb{E}[e^{tZ^2}] = (1 - 2t\sigma^2)^{-1/2}$$

Then given  $b$ , since  $a^\top V b$  has distribution  $\mathcal{N}(0, \|Vb\|^2)$ , we see immediately that whenever

$$t \|Vb\|^2 < \frac{1}{4},$$

$$\mathbb{E}\left[e^{t(a^\top Vb)^2} \mid b\right] = (1 - 2t \|Vb\|^2)^{-1/2} \leq e^{2t \|Vb\|^2}$$

where the latter bound comes from, if  $x \leq \frac{1}{4}$ ,

$$\log\left((1 - 2x)^{-1/2}\right) = -\frac{1}{2} \log(1 - 2x) \leq \frac{1}{2} \cdot \frac{2x}{1 - 2x} \leq 2x$$

While it is now tempting to use  $\|Vb\| \leq \|V\|_F \cdot \|b\|$ , this leads to a bad bound. Instead, observe that  $Vb$  is  $\mathcal{N}(0, VV^\top)$  distributed.

Now recall that for  $Z \sim \mathcal{N}(0, A)$ , if  $t < 1/(2\lambda_{\max}(A))$ ,

$$\mathbb{E}\left[e^{t\|Z\|^2}\right] = \det(I - 2tA)^{-1/2}$$

Therefore, as long as  $t < 1/(8\sigma_{\max}(V))$ , we have that

$$\begin{aligned} \mathbb{E}\left[e^{2t\|Vb\|^2}\right] &= \prod_{i=1}^r (1 - 4t\sigma_i^2)^{-1/2} = \exp\left(\sum_{i=1}^r -\frac{1}{2} \log(1 - 4t\sigma_i^2)\right) \\ &\leq \exp\left(4t \sum_{i=1}^r \sigma_i^2\right) = \exp\left(4t \|V\|_F^2\right) \end{aligned}$$

We are now almost done. Observe, on the event  $B$ ,

$$\|Vb\|^2 = b^\top V^\top Vb = \sum_{i=1}^r \sigma_i^2 b^\top q_i \leq 3 \log(C/\delta) \sum_{i=1}^r \sigma_i^2 = 3 \log(C/\delta) \|V\|_F^2$$

Now,  $t \|Vb\|^2 \leq t 3 \|V\|_F^2 \log(C/\delta)$ , so choosing  $t = 1/(24 \|V\|_F^2 \log(C/\delta))$ , we will satisfy all the bounds above ( $\sigma_{\max}(V) \leq \|V\|_F$ ), and get

$$\mathbb{E}\left[e^{t(a^\top Vb)^2}\right] = \mathbb{E}\left[\mathbb{E}\left[e^{t(a^\top Vb)^2} \mid b\right]\right] \leq e^{4t\|V\|_F^2} \leq \exp\left(\frac{1}{6 \log(C/\delta)}\right) \leq e^{1/6} \leq 2$$

Therefore,  $a_t^\top V_t b_t$  is  $\text{nSG}(\sqrt{24 \log(C/\delta)} \|V\|_F)$ . Combining the previous bounds on  $\|a\|$  and  $\|b\|$ , we get

$$\mathbb{E}_{\mathcal{F}_{t-1}} \left[ \exp\left(\frac{\|a_t a_t^\top V_t b_t b_t^\top \mathbb{1}_{A_t \cap B_t}\|_F^2}{216md \log^3(C/\delta) \|V\|_F^2}\right) \right] \leq 2$$

Therefore, using  $\|Q_t\|_F^2 \leq \|V_t\|^2 + \|a_t a_t^\top V_t b_t b_t^\top \mathbb{1}_{A_t \cap B_t}\|$ , for some absolute constant  $c'$ ,

$$\mathbb{E}_{\mathcal{F}_{t-1}} \left[ \exp \left( \frac{\|Q_t\|_F^2}{c' m d \log^3(C/\delta) \|V\|_F^2} \right) \right] \leq 2$$

This proves that  $Q_t$  is nSG( $c' \log^{3/2}(C/\delta) \|V\|_F$ ).

Thus, by Lemma 1, for some absolute constant  $c$ , with probability at least  $1 - \delta$ ,

$$\left\| \sum_{t=0}^{\tau-1} Q_t \right\| \leq c \cdot \theta \sum_{t=0}^{\tau-1} m d \log^3(C/\delta) \|V_t\|_F^2 + \frac{1}{\theta} \log \frac{2mn}{\delta}.$$

Now, on  $\bigcap_{t=0}^{\tau-1} A_t \cap B_t$ ,

$$\sum_{t=0}^{\tau-1} Q_t = \sum_{t=0}^{\tau-1} (V_t - a_t a_t^\top V_t b_t b_t^\top)$$

As  $\mathbb{P}_{\mathcal{F}_{t-1}}(B_{t,i}) \geq 1 - \delta$ , by union bound  $\mathbb{P}_{\mathcal{F}_{t-1}}(B_t) \geq 1 - m\delta$ . Further, recall that  $\mathbb{P}_{\mathcal{F}_{t-1}}(A_t) \geq 1 - \delta$ . Therefore,

$$\mathbb{P}_{\mathcal{F}_{t-1}} \left( \bigcap_{t=0}^{\tau-1} A_t \cap B_t \right) \geq 1 - \tau(m+1)\delta$$

so, with probability at least  $1 - \delta - \tau(m+1)\delta$ ,

$$\|E_\tau\| \leq c \cdot \theta \sum_{t=0}^{\tau-1} m d \log^3(C/\delta) \|V_t\|_F^2 + \frac{1}{\theta} \log \left( \frac{2md}{\delta} \right).$$

After rescaling  $\delta$  to  $\delta/(1 + \tau(m+1))$ , we have arrived at the desired result.  $\square$

We will need this definition to bound the quadratic error.

**Definition 5.9.** Random variable  $X \in \mathbb{R}$  is SW( $K, \alpha$ ) for some  $K, \alpha > 0$  if

$$\mathbb{P}(|X| \geq s) \leq 2 \exp \left( -(s/K)^{1/\alpha} \right) \quad \forall s \geq 0$$

We will need the following lemmas.

**Lemma 5.10.** Suppose  $X$  and  $Y$  are SW( $K_X, \alpha$ ) and SW( $K_Y, \alpha$ ) respectively. Then  $XY$  is SW( $C(K_X \cdot K_Y), 2\alpha$ ) and  $X + Y$  is SW( $C(K_X + K_Y), \alpha$ ) for some absolute constant  $C$ .

**Lemma 5.11.** *Suppose  $X_1, \dots, X_n$  are identically distributed  $SW(K', \alpha)$ . Then for some absolute constant  $c > 0$ , for all  $s \geq ncK'$ , we have*

$$\mathbb{P}\left(\left|\sum_{i=1}^n X_i\right| \geq s\right) \leq \exp\left(-\left(\frac{s}{ncK'}\right)^{1/\alpha}\right).$$

We will need to frequently bound sums of  $\|P_t\|^{2k}$ .

**Lemma 5.12.** *Suppose  $a_0, \dots, a_{\tau-1} \sim \mathcal{N}(0, I_d)$ , and  $b_0, \dots, b_{\tau-1} \sim \mathcal{N}(0, I_m)$ . Then for any  $k > 0$  there exists absolute constants  $c, C > 0$  such that for any  $0 < \delta < 1$ , w.p.  $\geq 1 - \delta$ ,*

$$\sum_{t=0}^{\tau-1} \|a_i b_i^\top\|_F^{2k} \leq \tau C c^{2k} d^k m^k (1 + \log(1/\delta)^{2k})$$

When  $\delta < 1/e$ ,

$$\sum_{t=0}^{\tau-1} \|a_i b_i^\top\|_F^{2k} \leq 2\tau C c^{2k} d^k m^k \log(1/\delta)^{2k}.$$

*Proof.*  $(a_i)_k^2$  and  $(b_j)_l^2$  are  $SW(1, 1)$ . Thus their product is  $SW(c, 2)$ . Thus  $\|a_i b_i^\top\|_F^2 = \sum_{k,l} (a_i)_k^2 (b_i)_l^2$  is  $SW(cmd, 2)$  for some absolute constant  $c$ . Thus,  $\|a_i b_i^\top\|_F^{2k}$  is  $SW(c^k m^k d^k, 2k)$ . Therefore by the lemma, for  $s \geq \tau C c^k m^k d^k$ ,

$$\mathbb{P}\left(\sum_{t=0}^{\tau-1} \|a_i b_i^\top\|_F^{2k} \geq s\right) \leq \exp\left(-\left(\frac{s}{\tau C c^k m^k d^k}\right)^{1/2k}\right)$$

the lemma follows by choosing  $s = (1 + \log(1/\delta)^{2k}) \tau C c^k m^k d^k$ .  $\square$

### 5.3 Function Decrease

First, we prove the following crucial lemma.

**Lemma 5.13** (Function decrease for rank  $r$  symmetric EGGROLL). *Suppose at each time  $t$ , the algorithm performs the update step with rank  $r \leq dm$ ,*

$$X_{t+1} = X_t - \eta g_u^r(X_t)$$

where

$$g_\sigma^r(X_t) := \frac{1}{r} \sum_{i=1}^r \frac{f(X_t + \sigma P_{t,i}) - f(X_t - \sigma P_{t,i})}{2\sigma} P_{t,i}$$

where  $P_{t,i} = a_{t,i} b_{t,i}^\top$  where each  $a_{t,i}, b_{t,i}$  are drawn i.i.d. from  $\mathcal{N}(0, I_d)$ , and  $\mathcal{N}(0, I_m)$  respectively, and  $\sigma > 0$  is the smoothing radius.

Then there exists absolute constants  $c > 0, C \geq 1$  such that, for any  $T \in \mathbb{Z}^+$  and  $1 \leq \tau' \leq T$ , and  $0 < \delta < 1/e$ , then upon defining  $\mathcal{H}_\tau(\delta)$  to be the event on which the inequality

$$\begin{aligned} f(X_\tau) - f(X_0) \leq & -\frac{\eta}{2} \sum_{t=0}^{\tau-1} \frac{1}{r} \sum_{i=1}^r |\langle \nabla f(X_t), P_{t,i} \rangle|^2 + \tau C c^4 \eta \sigma^2 L^2 (dm)^2 \log\left(\frac{T}{\delta}\right)^4 \\ & + \eta^2 \frac{cdm\chi^4}{r} \sum_{t=0}^{\tau-1} \|\nabla f(X_t)\|^2 + \eta^2 \tau C c^6 \sigma^2 L^3 (dm)^3 \log\left(\frac{T}{\delta}\right)^6 \end{aligned}$$

holds, (where  $\chi := \log(Cm^2 dr T / \delta)$ ), we have

$$\mathbb{P}\left(\bigcap_{\tau=1}^{\tau'} \mathcal{H}_\tau(\delta)\right) \geq 1 - \frac{3\tau'}{T} \delta$$

*Proof.* For each  $t \in \{-1, \dots, \tau\}$ , let

$$\mathcal{F}_t = \sigma(X_0, \{a_{0,i}\}_{i=1}^r, \dots, \{a_{t,i}\}_{i=1}^r, \{b_{0,i}\}_{i=1}^r, \dots, \{b_{t,i}\}_{i=1}^r)$$

be the sigma-algebra containing the information available at time  $t$ , and as usual write

$$P_{t,i} = a_{t,i} b_{t,i}^\top.$$

Observe that, by mean value theorem, there exists some  $s_{t,i} \in [-1, 1]$  such that

$$g_\sigma(X_t) = \frac{1}{r} \sum_{i=1}^r \frac{f(X_t + \sigma P_{t,i}) - f(X_t - \sigma P_{t,i})}{2\sigma} P_{t,i} = \frac{1}{r} \sum_{i=1}^r \nabla f(X_t + \sigma s_{t,i} P_{t,i}) P_{t,i}$$

Then we have that

$$X_{t+1} - X_t = -\eta \cdot \frac{1}{r} \sum_{i=1}^r (\langle \nabla f(X_t), P_{t,i} \rangle P_{t,i} + \langle \nabla f(X_t + \sigma s_{t,i} P_{t,i}) - \nabla f(X_t), P_{t,i} \rangle P_{t,i})$$

$$\begin{aligned}
& f(X_{t+1}) - f(X_t) \\
& \leq \langle X_{t+1} - X_t, \nabla f(X_t) \rangle + \frac{L}{2} \|X_{t+1} - X_t\|^2 \\
& \leq -\eta \cdot \frac{1}{r} \sum_{i=1}^r |\langle \nabla f(X_t), P_{t,i} \rangle|^2 \\
& \quad - \eta \cdot \frac{1}{r} \sum_{i=1}^r \langle \nabla f(X_t), P_{t,i} \rangle \langle \nabla f(X_t + \sigma_{t,i} P_{t,i}) - \nabla f(X_t), P_{t,i} \rangle \\
& \quad + \frac{L\eta^2}{2} \left\| \frac{1}{r} \sum_{i=1}^r (\langle \nabla f(X_t), P_{t,i} \rangle P_{t,i} + \langle \nabla f(X_t + \sigma_{t,i} P_{t,i}) - \nabla f(X_t), P_{t,i} \rangle P_{t,i}) \right\|^2
\end{aligned}$$

Now notice, by using  $ab \leq (a^2 + b^2)/2$ ,

$$\begin{aligned}
& -\frac{\eta}{r} \sum_{i=1}^r \langle \nabla f(X_t), P_{t,i} \rangle \langle \nabla f(X_t + \sigma_{t,i} P_{t,i}) - \nabla f(X_t), P_{t,i} \rangle \\
& \leq \frac{\eta}{2r} \sum_{i=1}^r \left( |\langle \nabla f(X_t), P_{t,i} \rangle|^2 + \sigma^2 L^2 \|P_{t,i}\|^4 \right)
\end{aligned}$$

Also, by  $\|a + b\|^2 \leq 2\|a\|^2 + 2\|b\|^2$ ,

$$\begin{aligned}
& \left\| \frac{1}{r} \sum_{i=1}^r (\langle \nabla f(X_t), P_{t,i} \rangle P_{t,i} + \langle \nabla f(X_t + \sigma_{t,i} P_{t,i}) - \nabla f(X_t), P_{t,i} \rangle P_{t,i}) \right\|^2 \\
& \leq 2 \left\| \frac{1}{r} \sum_{i=1}^r \langle \nabla f(X_t), P_{t,i} \rangle P_{t,i} \right\|^2 + \frac{2}{r} \sum_{i=1}^r \|\langle \nabla f(X_t + \sigma_{t,i} P_{t,i}) - \nabla f(X_t), P_{t,i} \rangle P_{t,i}\|^2 \\
& \leq 2 \left\| \frac{1}{r} \sum_{i=1}^r \langle \nabla f(X_t), P_{t,i} \rangle P_{t,i} \right\|^2 + \frac{2\sigma^2 L^2}{r} \sum_{i=1}^r \|P_{t,i}\|^6
\end{aligned}$$

Therefore,

$$\begin{aligned}
f(X_{t+1}) - f(X_t) & \leq -\frac{\eta}{2r} \sum_{i=1}^r |\langle \nabla f(X_t), P_{t,i} \rangle|^2 + \frac{\eta\sigma^2 L^2}{2r} \sum_{i=1}^r \|P_{t,i}\|^4 \\
& \quad + L\eta^2 \left\| \frac{1}{r} \sum_{i=1}^r \langle \nabla f(X_t), P_{t,i} \rangle P_{t,i} \right\|^2 + \frac{\eta^2 L^3 \sigma^2}{r} \sum_{i=1}^r \|P_{t,i}\|^6
\end{aligned}$$

Now,

$$\left\| \frac{1}{r} \sum_{i=1}^r \langle \nabla f(X_t), P_{t,i} \rangle P_{t,i} \right\|^2 \leq 2 \left\| \frac{1}{r} \sum_{i=1}^r (\nabla f(X_t) - \langle \nabla f(X_t), P_{t,i} \rangle P_{t,i}) \right\|^2 + 2 \|\nabla f(X_t)\|^2$$

By our proposition, with probability  $1 - \delta$ ,

$$\left\| \sum_{i=1}^r (\nabla f(X_t) - \langle \nabla f(X_t), P_{t,i} \rangle P_{t,i}) \right\| \leq c\theta dmr \log^3 \left( \frac{Crm}{\delta} \right) \|\nabla f(X_t)\|^2 + \frac{1}{\theta} \log \left( \frac{Cm^2 dr}{\delta} \right)$$

As  $\nabla f(X_t)$  is fixed given  $\mathcal{F}_{t-1}$ , optimizing over  $\theta$ , this is at most

$$\sqrt{cdmr} \log(Cm^2 dr/\delta)^2 \|\nabla f(X_t)\|.$$

Therefore for some irrelevant absolute constant  $c$ ,

$$\left\| \frac{1}{r} \sum_{i=1}^r \langle \nabla f(X_t), P_{t,i} \rangle P_{t,i} \right\|^2 \leq c \frac{dm}{r} \log(Cm^2 dr/\delta)^4 \|\nabla f(X_t)\|^2 \quad (5.1)$$

By union bound, this will hold jointly for all  $0 \leq t \leq \tau' - 1$  with probability at least  $1 - \tau'\delta$ .

Now by old lemma, for some absolute constants  $C, c$  possibly different from before, with probability at least  $1 - \delta$ ,

$$\frac{1}{r} \sum_{t=0}^{\tau-1} \sum_{i=1}^r \|P_{t,i}\|^{2k} \leq 2\tau C c^{2k} (dm)^k \log(1/\delta)^{2k}.$$

If  $k$  is fixed, by union bound, this will hold jointly for all  $1 \leq \tau \leq \tau'$  with probability  $1 - \tau'\delta$ .

Therefore, it will hold for both  $k = 2$  and  $3$ , for all  $1 \leq \tau \leq \tau'$ , with probability at least  $1 - 2\tau'\delta$ .

Therefore,

$$\begin{aligned} f(X_\tau) - f(X_0) &\leq -\frac{\eta}{2r} \sum_{t=0}^{\tau-1} \sum_{i=1}^r |\langle \nabla f(X_t), P_{t,i} \rangle|^2 + \tau C c^4 \eta \sigma^2 L^2 (dm)^2 \log(1/\delta)^4 \\ &\quad + \eta^2 \frac{cdm}{r} \log(Cm^2 dr/\delta)^4 \sum_{t=0}^{\tau-1} \|\nabla f(X_t)\|^2 + \eta^2 \tau C c^6 \sigma^2 L^3 (dm)^3 \log(1/\delta)^6 \end{aligned}$$



holds for all  $1 \leq \tau \leq \tau'$  with probability at least  $1 - 3\tau\delta$ . Equivalently,

$$\mathbb{P}\left(\bigcap_{\tau=1}^{\tau'} \mathcal{H}_\tau(\delta)\right) \geq 1 - 3\tau'\delta$$

Similarly, we see that, as when  $\tau$  is fixed we need only Equation 5.1 to hold jointly over  $0 \leq t \leq \tau - 1$ , while the bound on the perturbation norms hold with probability at least  $1 - 2\delta$ , so,

$$\mathbb{P}(\mathcal{H}_\tau(\delta)) \geq 1 - (\tau + 2)\delta$$

After rescaling  $\delta$  to  $\delta/T$ , we have arrived at the desired result.  $\square$

## 5.4 Gradients Don't Move too Quickly

We will need the following Bernstein inequality for sub-exponential random variables (Vershynin, 2018, Theorem 2.8.1).

**Lemma 5.14** (Bernstein concentration inequality). *Let  $X_1, \dots, X_n$  be independent, mean-zero,  $\sigma$ -subexponential random variables, i.e.*

$$\mathbb{E}\left[\exp\left(\frac{|X_i|}{\sigma}\right)\right] \leq 2 \quad \text{for all } i \in [n].$$

*Then there exists an absolute constant  $c$  such that for any  $s \geq 0$ ,*

$$\mathbb{P}\left(\sum_{i=1}^n X_i \geq s\right) \leq \exp\left(-c \min\left\{\frac{s^2}{n\sigma^2}, \frac{s}{\sigma}\right\}\right).$$

We then prove the following lemma.

**Lemma 5.15.** *There exists an absolute constant  $c_2 \geq 1$  such that, if*

$$t_f(\delta) = \frac{c_2}{m} \log\left(\frac{T}{\delta}\right), \quad \delta > 0,$$

and

$$\mathcal{B}_{t_0}(\delta; k) = \bigcup_{t=t_0}^{t_0+k-1} \left\{ \frac{1}{m} \sum_{i=1}^m |a_{t,i}^\top \nabla f(X_t) b_{t,i}|^2 \geq \frac{1}{2} \|\nabla f(X_t)\|^2 \right\}$$

we have

$$\mathbb{P}(\mathcal{B}_{t_0}(\delta; k)) \geq 1 - \frac{\delta}{T}$$

for  $0 < \delta < 1$ ,  $t_0 \geq 0$  and  $k \geq t_f(\delta)$ .

*Proof.* By using the same technique from the proof of Theorem 3, we know that the set of random variables  $\left\{ \|\nabla f(X_t)\|^2 - |a_{t,i}^\top \nabla f(X_t) b_{t,i}|^2 \right\}_{i=1}^m$  are independent,  $c \|\nabla f(X_t)\|^2$  subexponential for some absolute constant  $c > 0$  given  $\mathcal{F}_{t-1}$ . Now define

$$E_t = \left\{ \frac{1}{m} \sum_{i=1}^m |a_{t,i}^\top \nabla f(X_t) b_{t,i}|^2 < \frac{1}{2} \|\nabla f(X_t)\|^2 \right\}$$

Therefore,

$$\begin{aligned} \mathbb{P}_{\mathcal{F}_{t-1}}(E_t) &= \mathbb{P}_{\mathcal{F}_{t-1}} \left( \frac{1}{m} \sum_{i=1}^m |a_{t,i}^\top \nabla f(X_t) b_{t,i}|^2 \geq \|\nabla f(X_t)\|^2 \right) \\ &= \mathbb{P}_{\mathcal{F}_{t-1}} \left( \sum_{i=1}^m \left( \|\nabla f(X_t)\|^2 - |a_{t,i}^\top \nabla f(X_t) b_{t,i}|^2 \right) > \frac{m}{2} \|\nabla f(X_t)\|^2 \right) \\ &\leq \exp \left( -c' \min \left\{ \frac{m}{4c^2}, \frac{m}{2c} \right\} \right) \leq \exp(-c''m) \end{aligned}$$

for some absolute constant  $c''$ , where the last line follows by using Bernstein concentration inequality.

Now,

$$\begin{aligned} \mathbb{P} \left( \bigcap_{t=t_0}^{t_0+k-1} E_t \right) &= \mathbb{P} \left( \bigcap_{t=t_0}^{t_0+k-1} \left\{ \frac{1}{m} \sum_{i=1}^m |a_{t,i}^\top \nabla f(X_t) b_{t,i}|^2 < \frac{1}{2} \|\nabla f(X_t)\|^2 \right\} \right) \\ &= \mathbb{E} \left[ \prod_{t=t_0}^{t_0+k-2} \mathbb{1}_{E_t} \mathbb{E}_{\mathcal{F}_{t_0+k-2}} [\mathbb{1}_{E_{t_0+k-1}}] \right] \leq \cdots \leq \exp(c''mk) \end{aligned}$$

therefore by choosing  $k \geq t_f(\delta)$  we get

$$\mathbb{P}\left(\bigcap_{t=t_0}^{t_0+k-1} \left\{ \frac{1}{m} \sum_{i=1}^m |a_{t,i} \nabla f(X_t) b_{t,i}|^2 < \frac{1}{2} \|\nabla f(X_t)\|^2 \right\}\right) \leq \frac{\delta}{T},$$

which completes the proof.  $\square$

**Lemma 5.16.** *Define*

$$\mathcal{A}_t(\delta) := \left\{ \|\nabla f(X_{t+1}) - \nabla f(X_t)\| \leq \frac{\|\nabla f(X_t)\|}{8t_f(\delta)} + \eta\sigma L^2 \cdot \frac{1}{r} \sum_{i=1}^r \|P_{t,i}\|^3 \right\}$$

where  $t_f(\delta)$  is defined above. Then there exists an absolute constant  $c_2$  such that, whenever  $\eta$  satisfies,

$$\eta L c_2 \sqrt{\frac{dm}{r}} \log\left(\frac{CdmrT}{\delta}\right)^2 \leq \frac{1}{8t_f(\delta)}, \quad (5.2)$$

we have

$$\mathbb{P}(\mathcal{A}_t(\delta)) \geq 1 - \frac{\delta}{T}$$

for any  $0 < \delta \leq 1/e$  and  $T \in \mathbb{Z}^+$ .

*Proof.* By  $L$ -lipschitz of  $\nabla f$ ,

$$\begin{aligned} \|\nabla f(X_{t+1}) - \nabla f(X_t)\| &\leq L \|X_{t+1} - X_t\| \\ &\leq \eta L \left\| \frac{1}{r} \sum_{i=1}^r a_{t,i} a_{t,i}^\top \nabla f(X_t) b_{t,i} b_{t,i}^\top \right\| \\ &\quad + \eta L \left\| \frac{1}{r} \sum_{i=1}^r \langle \nabla f(X_t + \sigma s_{t,i} P_{t,i}) - \nabla f(X_t), P_{t,i} \rangle P_{t,i} \right\| \end{aligned}$$

By  $L$ -lipschitz of  $\nabla f(X_t)$  and the Cauchy-Schwarz inequality,

$$\left\| \frac{1}{r} \sum_{i=1}^r \langle \nabla f(X_t + \sigma s_{t,i} P_{t,i}) - \nabla f(X_t), P_{t,i} \rangle P_{t,i} \right\| \leq L\sigma \cdot \frac{1}{r} \sum_{i=1}^r \|P_{t,i}\|^3$$

Now for some absolute constant  $c_2 > 0$  by Theorem 5.8, with probability at least  $1 - \delta/T$ ,

$$\left\| \frac{1}{r} \sum_{i=1}^r a_{t,i} a_{t,i}^\top \nabla f(X_t) b_{t,i} b_{t,i}^\top \right\| \leq c_2 \log(CdmrT/\delta)^2 \sqrt{\frac{dm}{r}} \|\nabla f(X_t)\|$$

plugging in our condition on  $\eta$ , we arrive at the statement of the lemma.  $\square$

## 5.5 Progress Lower Bound

This next very important lemma is the key that will let us lower bound

$$\sum_{t=0}^{\tau-1} \frac{1}{r} \sum_{i=1}^r |\langle \nabla f(X_t), P_t \rangle|^2 \geq \tilde{\Omega} \left( \sum_{t=0}^{\tau-1} \|\nabla f(X_t)\| \right)$$

in some sense.

Denote

$$\mathcal{E}(t_1, t_2, \delta) = \bigcap_{t=t_1}^{t_1+t_2-1} \left\{ \|\nabla f(X_t)\| \geq 8t_f(\delta)\eta\sigma L^2 \cdot \frac{1}{r} \sum_{i=1}^r \|P_{t,i}\|^3 \right\}$$

**Lemma 5.17.** *Let  $0 < \delta \leq 1/e$  and  $T \in \mathbb{Z}^+$  satisfying  $T \geq 2t_f(\delta) + 1$ . Consider any positive integer  $t'_f \leq 2t_f(\delta)$ , and any  $t_0 \in \{0, \dots, T - 1 - t'_f\}$ . Suppose  $\eta$  satisfies the condition above, then on*

$$\mathcal{E}(t_0, t'_f, \delta) \cap \left( \bigcap_{t=t_0}^{t_0+t'_f-1} \mathcal{A}_t(\delta) \right)$$

we have

$$\frac{1}{2} \|\nabla f(X_0)\| \leq \|\nabla f(X_t)\| \leq 2 \|\nabla f(X_t)\|$$

for all  $t \in \{t_0, \dots, t_0 + t'_f - 1\}$ .

*Proof.* By plugging in the definition for  $\mathcal{E}_t(\delta)$  into the definition of  $\mathcal{A}_t(\delta)$ , we see that on the prescribed event,

$$\|\nabla f(X_{t+1}) - \nabla f(X_t)\| \leq \frac{\|\nabla f(X_t)\|}{4t_f(\delta)}$$

therefore

$$\left(1 - \frac{1}{4t_f(\delta)}\right) \|\nabla f(X_t)\| \leq \|\nabla f(X_{t+1})\| \leq \left(1 + \frac{1}{4t_f(\delta)}\right) \|\nabla f(X_t)\|$$

which leads to

$$\left(1 - \frac{1}{4t_f(\delta)}\right)^{t-t_0} \|\nabla f(X_0)\| \leq \|\nabla f(X_t)\| \leq \left(1 + \frac{1}{4t_f(\delta)}\right)^{t-t_0} \|\nabla f(X_0)\|.$$

for all  $t \in \{t_0, \dots, t_0 + t'_f - 1\}$ . Now using  $\log(1+x) \leq x$  for  $x \geq 0$  and  $-\log(1-x) \leq x/(1-x)$  for  $0 \leq x < 1$ , we know that

$$\left(1 - \frac{1}{4x}\right)^{2x} \geq e^{-2/3} \geq \frac{1}{2}, \quad \left(1 + \frac{1}{4x}\right)^{2x} \leq e^{1/2} \leq 2$$

Now recalling that  $t'_f \leq 2t_f(\delta)$ , we are done.  $\square$

**Lemma 5.18.** *There exists an absolute constant  $c > 0$  such that for any  $t \in \mathbb{N}$ , the event*

$$\mathcal{G}_t(\delta) := \left\{ \frac{1}{r} \sum_{i=1}^r \|P_{t,i}\|^3 \leq c(dm)^{3/2} \log(T/\delta)^3 \right\}$$

*holds with probability at least  $1 - \delta/T$  for  $0 < \delta \leq 1/e$ .*

*Proof.* Obvious from Lemma 5.12 and the Cauchy-Schwarz inequality.  $\square$

**Lemma 5.19.** *Let  $0 < \delta \leq 1/e$  and  $T \in \mathbb{Z}^+$  such that  $T > 2t_f(\delta) + 1$ . Consider any positive integer  $t'_f \leq 2t_f(\delta)$  and  $t_0 \in \{0, \dots, T - 1 - t'_f\}$ . Suppose  $\eta$  satisfies the condition above. Then on the event*

$$\mathcal{E}^c(t_0, t'_f, \delta) \cap \left( \bigcap_{t=t_0}^{t_0+t'_f-1} \mathcal{A}_t(\delta) \right) \cap \left( \bigcap_{t=t_0}^{t_0+t'_f-1} \mathcal{G}_t(\delta) \right)$$

*we have*

$$\|\nabla f(X_t)\| \leq 16t_f(\delta)\eta\sigma L^2 c(dm)^{3/2} \log(T/\delta)^3$$

*for all  $t \in \{t_0, \dots, t_0 + t'_f - 1\}$ .*

*Proof.* Let  $t'$  be the first time in  $\{t_0, \dots, t_0 + t_f - 1\}$  such that the gradient is dominated, meaning

$$\|\nabla f(X_{t'})\| \leq 8t_f(\delta)\eta L^2\sigma \frac{1}{r} \sum_{i=1}^r \|P_{t,i}\|^3.$$

By definition of  $\mathcal{G}_{t'}(\delta)$ , we must have

$$\|\nabla f(X_{t'})\| \leq 8t_f(\delta)\eta L^2\sigma c(dm)^{3/2} \log(T/\delta)^3$$

For  $j < t'$ , using the idea from a previous lemma,

$$\|\nabla f(X_j)\| \leq 2\|\nabla f(X_{t'})\| \leq 16t_f(\delta)\eta L^2\sigma c(dm)^{3/2} \log(T/\delta)^3$$

On the other hand, for  $t \in \{t', \dots, t_0 + t'_f - 1\}$ , by using definition of  $\mathcal{A}_t(\delta)$  and  $\mathcal{G}_t(\delta)$ , we get

$$\begin{aligned} \|\nabla f(X_{t+1})\| &\leq \left(1 + \frac{1}{8t_f(\delta)}\right) \|\nabla f(X_t)\| + \eta\sigma L^2 c(dm)^{3/2} \log(T/\delta)^3 \\ &= \left(1 - \frac{1}{8t_f(\delta)}\right)^{t+1-t'} \|\nabla f(X_{t'})\| \\ &\quad + \sum_{i=0}^{t-t'} \left(1 - \frac{1}{8t_f(\delta)}\right)^{t-t'-i} \eta\sigma L^2 c(dm)^{3/2} \log(T/\delta)^3 \\ &\leq \left(1 + \frac{1}{8t_f(\delta)}\right)^{t'_f} \|\nabla f(X_{t'})\| \\ &\quad + 8t_f(\delta) \left( \left(1 + \frac{1}{8t_f(\delta)}\right)^{t'_f} - 1 \right) \eta\sigma L^2 c(dm)^{3/2} \log(T/\delta)^3 \\ &\leq e^{1/4} 8t_f(\delta) \eta\sigma L^2 c(dm)^{3/2} \log(T/\delta)^3 \\ &\quad + 8t_f(\delta) (e^{1/4} - 1) \eta\sigma L^2 c(dm)^{3/2} \log(T/\delta)^3 \\ &\leq 16t_f(\delta) \eta\sigma L^2 c(dm)^{3/2} \log(T/\delta)^3 \end{aligned}$$

Where we used  $t'_f \leq 2t_f(\delta)$  and  $(1 + 1/(8x))^{2x} \leq e^{1/4}$  when  $x > 0$ . This completes the proof.  $\square$

## 5.6 Putting it all Together

We are now almost ready to prove the main theorem. But first, we need the following lemma.

**Lemma 5.20** (Function change for large  $\tau$ ). *Let  $0 < \delta \leq 1/e$  and let  $\tau \geq t_f(\delta)$  be arbitrary. Split  $\{0, \dots, \tau - 1\}$  into  $K := \tau/t_f(\delta)$  intervals  $J_0, \dots, J_{K-1}$  of size  $t_f(\delta)$ . Let  $I_1$  denote the set of indices where the gradient always dominates the noise term in  $J_k$ , i.e.*

$$I_1 := \left\{ k : \|\nabla f(X_t)\| \geq 8t_f(\delta)\eta\sigma L^2 \frac{1}{r} \sum_{i=1}^r \|P_{t,i}\|^3 \text{ for all } t \in J_k \right\}$$

Suppose we choose  $\eta$  satisfying

$$\eta \leq \frac{1}{8t_f(\delta)} \min \left\{ \sqrt{\frac{r}{dm}} \cdot \frac{1}{c_2 L \log(Cdmr\tau/\delta)^2}, \frac{r}{dm} \cdot \frac{1}{8c\chi^4} \right\}$$

Then on the event

$$\mathcal{E}_\tau(\delta) := \mathcal{H}_\tau(\delta) \cap \left( \bigcap_{t=0}^{\tau-1} \mathcal{A}_t(\delta) \right) \cap \left( \bigcap_{t=0}^{\tau-1} \mathcal{G}_t(\delta) \right) \cap \left( \bigcap_{t=0}^{\tau-1} \mathcal{B}_{kt_f(\delta)}(\delta; t_f(\delta)) \right)$$

we have the following upper bound for some absolute constants  $c, C$ ,

$$\begin{aligned} f(X_\tau) - f(X_0) &\leq - \sum_{k \in I_1} \frac{\eta}{4} \min_{t \in J_k} \|\nabla f(X_t)\|^2 + \tau \eta^4 \frac{cdm\chi^4}{r} t_f(\delta)^2 \sigma^2 L^4 (dm)^3 \log(T/\delta)^6 \\ &\quad + \tau C c^4 \eta \sigma^2 L^2 (dm)^2 \log\left(\frac{T}{\delta}\right)^4 + \eta^2 \tau C c^6 \sigma^2 L^3 (dm)^3 \log\left(\frac{T}{\delta}\right)^6. \end{aligned}$$

Furthermore,

$$\mathbb{P}(\mathcal{E}_\tau(\delta)) \geq 1 - \frac{(4\tau + 2)}{T} \delta.$$

*Proof.* Any interval  $J_k = \{t_0, \dots, t_0 + t_f - 1\}$  belongs to one of the following two cases:

**(Case 1)** (Gradient dominates noise): This means for every  $t \in J_k$ ,

$$\|\nabla f(X_t)\| \geq 8t_f(\delta)\eta\sigma L^2 \frac{1}{r} \sum_{i=1}^r \|P_{t,i}\|^3$$

By the Lemma 5.17 we have

$$\min_{t \in J_k} \|\nabla f(X_t)\| \geq \frac{1}{4} \max_{t \in J_k} \|\nabla f(X_t)\|$$

On the event  $\mathcal{B}_{kt_f(\delta)}(\delta; t_f(\delta))$ , there exists some  $t \in J_k$  such that

$$\frac{1}{r} \sum_{i=1}^r |\langle P_{t,i}, \nabla f(X_t) \rangle|^2 \geq \frac{1}{2} \|\nabla f(X_t)\|^2$$

Therefore, using an average,

$$\begin{aligned} \frac{1}{4} \sum_{t \in J_k} \frac{1}{r} \sum_{i=1}^r |\langle P_{t,i}, \nabla f(X_t) \rangle|^2 &\geq \frac{1}{4} \min_{t \in J_k} \|\nabla f(X_t)\| \geq \frac{1}{64} \max_{t \in J_k} \|\nabla f(X_t)\|^2 \\ &\geq \frac{1}{64t_f(\delta)} \sum_{t \in J_k} \|\nabla f(X_t)\|^2 \end{aligned}$$

By choosing

$$\frac{\eta c d m \chi^4}{r} \leq \frac{1}{64t_f(\delta)}$$

we have

$$\begin{aligned} & -\frac{\eta}{2} \sum_{t \in J_k} \frac{1}{r} \sum_{i=1}^r |\langle P_{t,i}, \nabla f(X_t) \rangle|^2 + \frac{\eta^2 c d m \chi^4}{r} \sum_{t \in J_k} \|\nabla f(X_t)\|^2 \\ & \leq -\frac{\eta}{4} \sum_{t \in J_k} \frac{1}{r} \sum_{i=1}^r |\langle P_{t,i}, \nabla f(X_t) \rangle|^2 \leq -\frac{\eta}{4} \min_{t \in J_k} \|\nabla f(X_t)\|^2. \end{aligned}$$

**(Case 2)** (Gradient does not dominate noise): This means that there exists some  $t \in J_k$  such that

$$\|\nabla f(X_t)\| \leq 8t_f(\delta)\eta\sigma L^2 \frac{1}{r} \sum_{i=1}^r \|P_{t,i}\|^3$$

Then by Lemma 5.19, we know that, for all  $t \in J_k$ ,

$$\|\nabla f(X_t)\| \leq 16t_f(\delta)\eta\sigma L^2 c(dm)^{3/2} \log(T/\delta)^3$$



Every  $J_k$  falls into one of the above two cases. Now we can complete the proof.

$$\begin{aligned}
& -\frac{\eta}{2} \sum_{t=0}^{\tau-1} \frac{1}{r} \sum_{i=1}^r |\langle P_{t,i}, \nabla f(X_t) \rangle|^2 + \eta^2 \frac{cdm\chi^4}{r} \sum_{t=0}^{\tau-1} \|\nabla f(X_t)\|^2 \\
& = \sum_{k \in I_1} \left( -\frac{\eta}{2} \sum_{t \in J_k} \frac{1}{r} \sum_{i=1}^r |\langle P_{t,i}, \nabla f(X_t) \rangle|^2 + \eta^2 \frac{cdm\chi^4}{r} \sum_{t \in J_k} \|\nabla f(X_t)\|^2 \right) \\
& \quad + \sum_{k \in I_1^c} \left( -\frac{\eta}{2} \sum_{t \in J_k} \frac{1}{r} \sum_{i=1}^r |\langle P_{t,i}, \nabla f(X_t) \rangle|^2 + \eta^2 \frac{cdm\chi^4}{r} \sum_{t \in J_k} \|\nabla f(X_t)\|^2 \right) \\
& \leq - \sum_{k \in I_1} \frac{\eta}{4} \min_{t \in J_k} \|\nabla f(X_t)\|^2 + \sum_{k \in I_1^c} t_f(\delta) \left( \eta^4 \frac{cdm\chi^4}{r} 256 t_f(\delta)^2 \sigma^2 L^4 c(dm)^3 \log(T/\delta)^6 \right) \\
& \leq - \sum_{k \in I_1} \frac{\eta}{4} \min_{t \in J_k} \|\nabla f(X_t)\|^2 + \tau \eta^4 \frac{cdm\chi^4}{r} 256 t_f(\delta)^2 \sigma^2 L^4 c(dm)^3 \log(T/\delta)^6
\end{aligned}$$

and thus by the definition of  $\mathcal{H}_\tau(\delta)$  in Lemma 5.13,

$$\begin{aligned}
f(X_\tau) - f(X_0) & \leq - \sum_{k \in I_1} \frac{\eta}{4} \min_{t \in J_k} \|\nabla f(X_t)\|^2 + \tau \eta^4 \frac{cdm\chi^4}{r} 256 t_f(\delta)^2 \sigma^2 L^4 c(dm)^3 \log(T/\delta)^6 \\
& \quad + \tau C c^4 \eta \sigma^2 L^2 (dm)^2 \log\left(\frac{T}{\delta}\right)^4 + \eta^2 \tau C c^6 \sigma^2 L^3 (dm)^3 \log\left(\frac{T}{\delta}\right)^6.
\end{aligned}$$

Furthermore,

$$\begin{aligned}
\mathbb{P}(\mathcal{E}_\tau(\delta)^c) & \leq \mathbb{P}(\mathcal{H}_\tau(\delta)^c) + \sum_{t=0}^{\tau-1} \mathbb{P}(\mathcal{A}_t(\delta)^c) + \sum_{t=0}^{\tau-1} \mathbb{P}(\mathcal{G}_t(\delta)^c) + \sum_{k=0}^{K-1} \mathbb{P}(\mathcal{B}_{kt_f(\delta)}(\delta; t_f(\delta))^c) \\
& \leq \frac{(\tau+2)}{T} \delta + \frac{\tau}{T} \delta + \frac{\tau}{T} \delta + \frac{K}{T} \delta \leq \frac{(4\tau+2)}{T} \delta
\end{aligned}$$

which completes the proof.  $\square$

**Theorem 5.21** (Symmetric EGGROLL Convergence). *Let  $0 < \delta < 1/e$  be arbitrary.*

Suppose we choose  $\sigma, \eta$  and  $T$  such that

$$\begin{aligned}\sigma^2 &\leq \min \left\{ \frac{\varepsilon^2}{1024Cc^4L^2(dm)^2 \log(T/\delta)^4}, \frac{\varepsilon^2}{8096rL^3(dm)^2 \log(T/\delta)^2} \right\} \\ \eta &\leq \frac{1}{8t_f(\delta)} \min \left\{ \sqrt{\frac{r}{dm}} \cdot \frac{1}{cL \log(Cdmr\tau/\delta)^2}, \frac{r}{dm} \cdot \frac{1}{8c\chi^4} \right\}, \\ T &\geq \max \left\{ 4, \frac{512t_f(\delta)(f(X_0) - f^*)}{\eta\varepsilon^2} \right\}.\end{aligned}$$

Then with probability at least  $1 - 5\delta$  there are at most  $T/4$  iterations for which  $\|\nabla f(X_t)\| \geq \varepsilon$ .

*Proof.* Without loss of generality assume that  $t_f(\delta)$  divides  $T$  and split  $\{0, \dots, T-1\}$  into  $K = T/t_f(\delta)$  intervals  $J_0, \dots, J_{K-1}$ . Let  $I_1$  denote the set of indices  $k$  such that for every  $t \in J_k$ ,

$$\|\nabla f(X_t)\| \geq 8t_f(\delta)\eta\sigma L^2 \frac{1}{r} \sum_{i=1}^r \|P_{t,i}\|^3.$$

Denote

$$\mathcal{E}_T(\delta) := \mathcal{H}_T(\delta) \cap \left( \bigcap_{t=0}^{T-1} \mathcal{A}_t(\delta) \right) \cap \left( \bigcap_{t=0}^{T-1} \mathcal{G}_t(\delta) \right) \cap \left( \bigcap_{k=0}^{K-1} \mathcal{B}_{kt_f(\delta)}(\delta; t_f(\delta)) \right)$$

For the rest of the proof assume we are working on  $\mathcal{E}_T(\delta)$ . By Lemma 5.20, and our choices of  $\eta$  and  $\delta$  in the statement of the theorem, we have

$$\begin{aligned}f(X_T) - f(X_0) &\leq - \sum_{k \in I_1} \frac{\eta}{4} \min_{t \in J_k} \|\nabla f(X_t)\|^2 + T\eta^4 \frac{cdm\chi^4}{r} t_f(\delta)^2 \sigma^2 L^4 (dm)^3 \log(T/\delta)^6 \\ &\quad + T C c^4 \eta \sigma^2 L^2 (dm)^2 \log\left(\frac{T}{\delta}\right)^4 + \eta^2 T C c^6 \sigma^2 L^3 (dm)^3 \log\left(\frac{T}{\delta}\right)^6\end{aligned}$$

Assume that there are more than  $T/4$  iterations with  $\|\nabla f(X_t)\| \geq \varepsilon$ . Let  $I_\varepsilon$  denote the set of indices  $k$  where there is a  $t \in J_k$  with  $\|\nabla f(X_t)\| \geq \varepsilon$ . Thus by pigeonhole principle,  $|I_\varepsilon| \geq T/4t_f(\delta)$ . By our choices of parameters, it can be shown that

$$16t_f(\delta)\eta\sigma L^2 c(dm)^{3/2} \log(T/\delta)^3 < \varepsilon$$

While by Lemma 5.19, if  $k \in I_1^c$ ,

$$\|\nabla f(X_t)\| \leq 16t_f(\delta)\eta\sigma L^2 c(dm)^{3/2} \log(T/\delta)^3$$

Therefore,  $I_\varepsilon \subset I_1$ . Now by Lemma 5.17, for any  $k \in I_1$ ,

$$\frac{1}{2} \left\| \nabla f(X_{kt_f(\delta)}) \right\| \leq \left\| \nabla f(X_t) \right\| \leq 2 \left\| \nabla f(X_{kt_f(\delta)}) \right\|, \quad \forall t \in J_k$$

This implies that for any  $k \in I_\varepsilon$ ,  $\min_{t \in J_k} \|\nabla f(X_t)\|^2 \geq \frac{1}{16}\varepsilon^2$ , and consequently,

$$-\sum_{k \in I_1} \frac{\eta}{4} \min_{t \in J_k} \|\nabla f(X_t)\|^2 \leq -\sum_{k \in I_\varepsilon} \frac{\eta}{2} \cdot \frac{\varepsilon^2}{16} \leq -\frac{T}{4t_f(\delta)} \cdot \frac{\eta}{4} \cdot \frac{\varepsilon^2}{16} = -\frac{T\eta\varepsilon^2}{256t_f(\delta)}$$

Hence,

$$\begin{aligned} f(X_T) - f(X_0) &\leq -\frac{T\eta\varepsilon^2}{256t_f(\delta)} + T\eta^4 \frac{cdm\chi^4}{r} t_f(\delta)^2 \sigma^2 L^4(dm)^3 \log(T/\delta)^6 \\ &\quad + TCc^4\eta\sigma^2 L^2(dm)^2 \log\left(\frac{T}{\delta}\right)^4 + \eta^2 TCc^6\sigma^2 L^3(dm)^3 \log\left(\frac{T}{\delta}\right)^6 \end{aligned}$$

By our choice of parameters,

$$\begin{aligned} T\eta^4 \frac{cdm\chi^4}{r} t_f(\delta)^2 \sigma^2 L^4(dm)^3 \log(T/\delta)^6 + TCc^4\eta\sigma^2 L^2(dm)^2 \log\left(\frac{T}{\delta}\right)^4 \\ + \eta^2 TCc^6\sigma^2 L^3(dm)^3 \log\left(\frac{T}{\delta}\right)^6 \leq \frac{T\eta\varepsilon^2}{512t_f(\delta)} \end{aligned}$$

We thus get

$$f(X_T) - f(X_0) \leq -\frac{T\eta\varepsilon^2}{512t_f(\delta)}$$

Therefore, as long as

$$T \geq \frac{512t_f(\delta)(f(x_0) - f^*)}{\eta\varepsilon^2}$$

we would get  $f(X_T) < f^*$ , a contradiction. Thus, we may conclude that on the event  $\mathcal{E}_T(\delta)$  there are at most  $T/4$  iterations for which  $\|\nabla f(X_t)\| \geq \varepsilon$ .

Finally,

$$\begin{aligned}\mathbb{P}(\mathcal{E}_T^c) &\leq \mathbb{P}(\mathcal{H}_T^c) + \sum_{t=0}^{T-1} \mathbb{P}(\mathcal{A}_t(\delta)) + \sum_{t=0}^{T-1} \mathbb{P}(\mathcal{G}_t(\delta)) + \sum_{k=0}^{K-1} \mathbb{P}(\mathcal{B}_{kt_f(\delta)}(\delta; t_f(\delta))) \\ &\leq \frac{(T+2)\delta}{T} + \delta + \delta + \frac{K\delta}{T} \leq 5\delta.\end{aligned}$$

□

# Bibliography

- Amir Ali Ahmadi, Alex Olshevsky, Pablo A. Parrilo, and John N. Tsitsiklis. Np-hardness of deciding convexity of quartic polynomials and related problems. *Mathematical Programming*, 137(1–2):453–476, 2013. doi: 10.1007/s10107-011-0499-2.
- Zeyuan Allen-Zhu and Lorenzo Orecchia. Linear coupling: An ultimate unification of gradient and mirror descent. *arXiv preprint arXiv:1407.1537*, 2016. doi: 10.48550/arXiv.1407.1537. URL <https://arxiv.org/abs/1407.1537>.
- Animashree Anandkumar and Rong Ge. Efficient approaches for escaping higher order saddle points in non-convex optimization. In *Proceedings of the 29th Annual Conference on Learning Theory*, volume 49 of *Proceedings of Machine Learning Research*, pages 81–102. PMLR, 2016. URL <https://proceedings.mlr.press/v49/anandkumar16.html>.
- Krishnakumar Balasubramanian and Saeed Ghadimi. Zeroth-order (non)-convex stochastic optimization via conditional gradient and gradient updates. In *Advances in Neural Information Processing Systems*, volume 31, 2018.
- Saeed Ghadimi and Guanghui Lan. Stochastic first- and zeroth-order methods for nonconvex stochastic programming. *arXiv preprint arXiv:1309.5549*, 2013. URL <https://arxiv.org/abs/1309.5549>.
- Nikolaus Hansen and Andreas Ostermeier. Completely derandomized self-adaptation in evolution strategies. *Evolutionary Computation*, 9(2):159–195, 2001.
- Chi Jin, Praneeth Netrapalli, Rong Ge, Sham M. Kakade, and Michael I. Jordan. A short note on concentration inequalities for random vectors with subgaussian norm. *arXiv preprint arXiv:1902.03736*, 2019. URL <https://arxiv.org/abs/1902.03736>.
- Béatrice Laurent and Pascal Massart. Adaptive estimation of a quadratic functional by model selection. *The Annals of Statistics*, 28(5):1302–1338, 2000. doi: 10.1214/aos/1015957395.
- Yurii Nesterov and Vladimir Spokoiny. Random gradient-free minimization of convex functions. *Foundations of Computational Mathematics*, 17(2):527–566, 2017. doi: 10.1007/s10208-015-9296-2.

- Yurii E. Nesterov. A method for solving the convex programming problem with convergence rate  $o(1/k^2)$ . *Soviet Mathematics Doklady*, 27(2):372–376, 1983.
- Ingo Rechenberg. Cybernetic solution path of an experimental problem. Technical Report Library Translation 1122, Royal Aircraft Establishment, Farnborough, UK, 1965. Translated by B. F. Toms.
- Ingo Rechenberg. *Evolutionstrategie: Optimierung technischer Systeme nach Prinzipien der biologischen Evolution*. Frommann-Holzboog, Stuttgart, 1973.
- Zhaolin Ren, Yujie Tang, and Na Li. Escaping saddle points in zeroth-order optimization: The power of two-point estimators. In *Proceedings of the 40th International Conference on Machine Learning*, 2023. URL <https://arxiv.org/abs/2209.13555>.
- Tim Salimans, Jonathan Ho, Xi Chen, Szymon Sidor, and Ilya Sutskever. Evolution strategies as a scalable alternative to reinforcement learning. *arXiv preprint arXiv:1703.03864*, 2017.
- Bidipta Sarkar, Mattie Fellows, Juan Agustin Duque, Alistair Letcher, Antonio León Villares, Anya Sims, Clarisse Wibault, Dmitry Samsonov, Dylan Cope, Jarek Liesen, Kang Li, Lukas Seier, Theo Wolf, Uljad Berdica, Valentin Mohl, Alexander David Goldie, Aaron Courville, Karin Sevegnani, Shimon Whiteson, and Jakob Nicolaus Foerster. Evolution strategies at the hyperscale. *arXiv preprint arXiv:2511.16652*, 2025. doi: 10.48550/arXiv.2511.16652. URL <https://arxiv.org/abs/2511.16652>.
- Weijie Su, Stephen Boyd, and Emmanuel J. Candès. A differential equation for modeling nesterov’s accelerated gradient method: Theory and insights. *Journal of Machine Learning Research*, 17(153):1–43, 2016. URL <https://www.jmlr.org/papers/v17/15-084.html>.
- Roman Vershynin. *High-Dimensional Probability: An Introduction with Applications in Data Science*. Cambridge University Press, 2018.
- Daan Wierstra, Tom Schaul, Tobias Glasmachers, Yi Sun, Jan Peters, and Jürgen Schmidhuber. Natural evolution strategies. *Journal of Machine Learning Research*, 15(27): 949–980, 2014.